

Inventory Management Dataset

According to the UN Environment Programme (UNEP), we waste one billion tons of food annually. Food waste from restaurants accounts for about 15% of all the food that ends up in landfills. This is an issue that has significant economic, human, and planetary consequences. One of the main reasons why restaurants produce so much food waste is because of improper inventory management, which leads to excess inventory that gets thrown away.

Using an ML model to analyze the history of past sales to predict the number of sales of each menu item can help restaurants optimize their inventory. With a better approximation of expected sales, restaurants can avoid excess inventory and food waste.

Dataset

Link - <https://data.mendeley.com/datasets/7s3ys9ntnz/>

The dataset that we are analyzing contains over two years of daily sales from a restaurant in Stuttgart, Germany. Number of servings sold for 7 products (CALAMARI, FISCH, GARNELEN, HAEHNCHEN, KOEFTE, LAMM, STEAK). Features include calendar informations such as weekday or special events, information derived from historical sales as well as weather from the local weather station.

ML Algorithms

1. Poisson Regression
 - Our ML output will be a count (number of sales for each menu item), so Poisson Regression will be a good fit in this case.
2. Tree Based Model
 - We want to capture complex non linear relationships between historical sales, weather, and menu sales. Our dataset contains many dimensions, so a tree based model would be a good fit to capture these complex interactions.
3. Time series
 - The data is contained as a sequence that was collected within the span of 2 years. This would be an ideal model for our dataset because we are predicting future sales based on past observations of sales, weather patterns, and seasonal patterns

Data Preparation

In [495...]

```
import pandas as pd
import seaborn as sns
import numpy as np
import warnings
import matplotlib.pyplot as plt
from sklearn.feature_selection import VarianceThreshold
from sklearn.linear_model import PoissonRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint
from sklearn.preprocessing import StandardScaler
from xgboost import XGBRegressor
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.statespace.sarimax import SARIMAX
from sklearn.feature_selection import RFE
from sklearn.model_selection import TimeSeriesSplit
from statsmodels.tsa.statespace.varmax import VARMAX
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
warnings.filterwarnings('ignore')
```

In [496...]

```
pd.set_option('display.max_columns', None)

df = pd.read_csv('data_preprocessed.csv', delimiter=';', encoding='utf-8')
df.head()
```

Out[496]:

	Unnamed: 0	DEMAND_DATE	Montag	Dienstag	Mittwoch	Donnerstag	Freitag	Samstag	Sonntag	MONTH_JAN	MONTH_FEB	MONTH_MAR
0	1	10.4.2013	0	0	0	0	1	0	0	0	0	0
1	2	10.5.2013	0	0	0	0	0	1	0	0	0	0
2	3	10.6.2013	0	0	0	0	0	0	1	0	0	0
3	4	10.7.2013	1	0	0	0	0	0	0	0	0	0
4	5	10.8.2013	0	1	0	0	0	0	0	0	0	0

In [497]: df.shape

Out[497]: (760, 293)

In [498]: df.describe()

	Unnamed: 0	Montag	Dienstag	Mittwoch	Donnerstag	Freitag	Samstag	Sonntag	MONTH_JAN	MONTH_
count	760.000000	760.000000	760.000000	760.000000	760.000000	760.000000	760.000000	760.000000	760.000000	760.000000
mean	380.500000	0.143421	0.143421	0.139474	0.142105	0.143421	0.144737	0.143421	0.081579	0.073
std	219.537392	0.350733	0.350733	0.346668	0.349388	0.350733	0.352067	0.350733	0.273902	0.261
min	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
25%	190.750000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
50%	380.500000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
75%	570.250000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
max	760.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000

In [499]: missing_values = df.isnull().sum()
print(missing_values[missing_values > 0])

Series([], dtype: int64)

In [500]: duplicate_rows = df[df.duplicated()]
print(f"Number of duplicate rows: {duplicate_rows.shape[0]}")

Number of duplicate rows: 0

In [501]: df = df.drop(columns=['Unnamed: 0', 'TOTAL_FISCHPROD', 'TOTAL_FLEISCH'])

df_time_series = df[['DEMAND_DATE', 'TOTAL']]
df = df.drop(columns=['DEMAND_DATE', 'TOTAL'])

df.head()

Out[501]:	Montag	Dienstag	Mittwoch	Donnerstag	Freitag	Samstag	Sonntag	MONTH_JAN	MONTH_FEB	MONTH_MAR	MONTH_APR	MONTH_MAY
0	0	0	0	0	1	0	0	0	0	0	0	0
1	0	0	0	0	0	1	0	0	0	0	0	0
2	0	0	0	0	0	0	1	0	0	0	0	0
3	1	0	0	0	0	0	0	0	0	0	0	0
4	0	1	0	0	0	0	0	0	0	0	0	0

In [489...]

#Dropping columns with low variance

```
selector = VarianceThreshold(threshold=0.01)
selector.fit(df)

low_variance_columns = df.columns[~selector.get_support()]
print(f"Low variance columns: {low_variance_columns}")

df = df.drop(columns=low_variance_columns)
```

Low variance columns: Index(['ISOUTLIER_HIGH', 'ISOUTLIER_LOW'], dtype='object')

In [490...]

```
# Drop highly correlated columns
correlation_matrix = df.corr().abs()
upper_tri = correlation_matrix.where(
    np.triu(np.ones(correlation_matrix.shape), k=1).astype(np.bool)
)

to_drop = [column for column in upper_tri.columns if any(upper_tri[column] > 0.8)]
df = df.drop(columns=to_drop)
print(f"Dropped columns due to high correlation: {to_drop}")
```

Dropped columns due to high correlation: ['TOTAL_FISCHPROD_DEMAND_T1', 'TOTAL_FISCHPROD_DEMAND_T2', 'TOTAL_FI SCHPROD_DEMAND_T3', 'TOTAL_FISCHPROD_DEMAND_T4', 'TOTAL_FISCHPROD_DEMAND_T5', 'TOTAL_FISCHPROD_DEMAND_T6', 'T OTAL_FISCHPROD_DEMAND_T7', 'TOTAL_FLEISCH_DEMAND_T1', 'TOTAL_FLEISCH_DEMAND_T2', 'TOTAL_FLEISCH_DEMAND_T3', 'TOTAL_FLEISCH_DEMAND_T4', 'TOTAL_FLEISCH_DEMAND_T5', 'TOTAL_FLEISCH_DEMAND_T6', 'TOTAL_FLEISCH_DEMAND_T7', 'TOTAL_DEMAND_T1', 'TOTAL_DEMAND_T2', 'TOTAL_DEMAND_T3', 'TOTAL_DEMAND_T4', 'TOTAL_DEMAND_T5', 'TOTAL_DEMAND_T6', 'TOTAL_DEMAND_T7', 'CALAMARI_CUM_DEMAND_T3', 'CALAMARI_CUM_DEMAND_T4', 'CALAMARI_CUM_DEMAND_T5', 'CALAMA RI_CUM_DEMAND_T6', 'CALAMARI_CUM_DEMAND_T7', 'FISCH_CUM_DEMAND_T3', 'FISCH_CUM_DEMAND_T4', 'FISCH_CUM_DEMAND_T5', 'FISCH_CUM_DEMAND_T6', 'FISCH_CUM_DEMAND_T7', 'GARNELEN_CUM_DEMAND_T3', 'GARNELEN_CUM_DEMAND_T4', 'GARNE LEN_CUM_DEMAND_T5', 'GARNELEN_CUM_DEMAND_T6', 'GARNELEN_CUM_DEMAND_T7', 'HAEHNCHEN_CUM_DEMAND_T3', 'HAEHNCHEN _CUM_DEMAND_T4', 'HAEHNCHEN_CUM_DEMAND_T5', 'HAEHNCHEN_CUM_DEMAND_T6', 'HAEHNCHEN_CUM_DEMAND_T7', 'KOEFT E_CUM_DEMAND_T3', 'KOEFT E_CUM_DEMAND_T4', 'KOEFT E_CUM_DEMAND_T5', 'KOEFT E_CUM_DEMAND_T6', 'KOEFT E_CUM_DEMAND_T7', 'LAMM_CUM_DEMAND_T3', 'LAMM_CUM_DEMAND_T4', 'LAMM_CUM_DEMAND_T5', 'LAMM_CUM_DEMAND_T6', 'LAMM_CUM_DEMAND_T7', 'STEAK_CUM_DEMAND_T3', 'STEAK_CUM_DEMAND_T4', 'STEAK_CUM_DEMAND_T5', 'STEAK_CUM_DEMAND_T6', 'STEAK_CUM_DEMAND_T7', 'TOTAL_FISCHPROD_CUM_DEMAND_T2', 'TOTAL_FISCHPROD_CUM_DEMAND_T3', 'TOTAL_FISCHPROD_CUM_DEMAND_T4', 'TOT AL_FISCHPROD_CUM_DEMAND_T5', 'TOTAL_FISCHPROD_CUM_DEMAND_T6', 'TOTAL_FISCHPROD_CUM_DEMAND_T7', 'TOTAL_FLEISCH _CUM_DEMAND_T2', 'TOTAL_FLEISCH_CUM_DEMAND_T3', 'TOTAL_FLEISCH_CUM_DEMAND_T4', 'TOTAL_FLEISCH_CUM_DEMAND_T5', 'TOTAL_FLEISCH_CUM_DEMAND_T6', 'TOTAL_FLEISCH_CUM_DEMAND_T7', 'TOTAL_CUM_DEMAND_T2', 'TOTAL_CUM_DEMAND_T3', 'TOTAL_CUM_DEMAND_T4', 'TOTAL_CUM_DEMAND_T5', 'TOTAL_CUM_DEMAND_T6', 'TOTAL_CUM_DEMAND_T7', 'TOTAL_HML_DEMAND _T7', 'TOTAL_NO_DAYS_ABOVE_7D_MEAN', 'GARNELEN_NO_DAYS_BELOW_7D_MEAN', 'HAEHNCHEN_NO_DAYS_BELOW_7D_MEAN', 'KO EFTE_NO_DAYS_BELOW_7D_MEAN', 'LAMM_NO_DAYS_BELOW_7D_MEAN', 'STEAK_NO_DAYS_BELOW_7D_MEAN', 'TOTAL_FISCHPROD_NO _DAYS_BELOW_7D_MEAN', 'TOTAL_FLEISCH_NO_DAYS_BELOW_7D_MEAN', 'TOTAL_NO_DAYS_BELOW_7D_MEAN', 'CALAMARI_MEAN_SA ME_WDAY_DEMANDS_W3', 'CALAMARI_MEAN_SAME_WDAY_DEMANDS_W4', 'FISCH_MEAN_SAME_WDAY_DEMANDS_W3', 'FISCH_MEAN_SAM E_WDAY_DEMANDS_W4', 'GARNELEN_MEAN_SAME_WDAY_DEMANDS_W2', 'GARNELEN_MEAN_SAME_WDAY_DEMANDS_W3', 'GARNELEN_MEA N_SAME_WDAY_DEMANDS_W4', 'HAEHNCHEN_MEAN_SAME_WDAY_DEMANDS_W2', 'HAEHNCHEN_MEAN_SAME_WDAY_DEMANDS_W3', 'HAEHN CHEN_MEAN_SAME_WDAY_DEMANDS_W4', 'KOEFT E_MEAN_SAME_WDAY_DEMANDS_W2', 'KOEFT E_MEAN_SAME_WDAY_DEMANDS_W3', 'KOE FTE_MEAN_SAME_WDAY_DEMANDS_W4', 'LAMM_MEAN_SAME_WDAY_DEMANDS_W2', 'LAMM_MEAN_SAME_WDAY_DEMANDS_W3', 'LAMM_MEA N_SAME_WDAY_DEMANDS_W4', 'STEAK_MEAN_SAME_WDAY_DEMANDS_W2', 'STEAK_MEAN_SAME_WDAY_DEMANDS_W3', 'STEAK_MEAN_SA ME_WDAY_DEMANDS_W4', 'TOTAL_FISCHPROD_MEAN_SAME_WDAY_DEMANDS_W2', 'TOTAL_FISCHPROD_MEAN_SAME_WDAY_DEMANDS_W 3', 'TOTAL_FISCHPROD_MEAN_SAME_WDAY_DEMANDS_W4', 'TOTAL_FLEISCH_MEAN_SAME_WDAY_DEMANDS_W2', 'TOTAL_FLEISCH_ME AN_SAME_WDAY_DEMANDS_W3', 'TOTAL_FLEISCH_MEAN_SAME_WDAY_DEMANDS_W4', 'TOTAL_MEAN_SAME_WDAY_DEMANDS_W2', 'TOT AL_MEAN_SAME_WDAY_DEMANDS_W3', 'TOTAL_MEAN_SAME_WDAY_DEMANDS_W4', 'SONNE', 'SONNE_T1', 'SONNE_T2', 'SONNE_T3', 'LUFTTEMPERATUR_T1', 'LUFTTEMPERATUR_T2', 'LUFTTEMPERATUR_T3', 'WIND_MEAN_OVER_LAST_T2', 'WIND_MEAN_OVER_LAST _T3', 'WIND_MEAN_OVER_LAST_T4', 'WIND_MEAN_OVER_LAST_T5', 'WIND_MEAN_OVER_LAST_T6', 'WIND_MEAN_OVER_LAST_T7', 'BEWOELKUNG_MEAN_OVER_LAST_T2', 'BEWOELKUNG_MEAN_OVER_LAST_T3', 'BEWOELKUNG_MEAN_OVER_LAST_T4', 'BEWOELKUNG_MEAN_OVER_LAST_T5', 'BEWOELKUNG_MEAN_OVER_LAST_T6', 'BEWOELKUNG_MEAN_OVER_LAST_T7', 'NIEDERSCHLAG_MEAN_OVER_LA ST_T3', 'NIEDERSCHLAG_MEAN_OVER_LAST_T4', 'NIEDERSCHLAG_MEAN_OVER_LAST_T5', 'NIEDERSCHLAG_MEAN_OVER_LAST_T6', 'NIEDERSCHLAG_MEAN_OVER_LAST_T7', 'SONNE_MEAN_OVER_LAST_T2', 'SONNE_MEAN_OVER_LAST_T3', 'SONNE_MEAN_OVER_LAST_T4', 'SONNE_MEAN_OVER_LAST_T5', 'SONNE_MEAN_OVER_LAST_T6', 'SONNE_MEAN_OVER_LAST_T7', 'LUFTTEMPERATUR_MEAN_O VER_LAST_T2', 'LUFTTEMPERATUR_MEAN_OVER_LAST_T3', 'LUFTTEMPERATUR_MEAN_OVER_LAST_T4', 'LUFTTEMPERATUR_MEAN_O VER_LAST_T5', 'LUFTTEMPERATUR_MEAN_OVER_LAST_T6', 'LUFTTEMPERATUR_MEAN_OVER_LAST_T7', 'NIEDERSCHLAG_HML_OVER_T 7', 'WIND_NO_DAYS_BELOW_7D_MEAN', 'BEWOELKUNG_NO_DAYS_BELOW_7D_MEAN', 'NIEDERSCHLAG_NO_DAYS_BELOW_7D_MEAN', 'SONNE_NO_DAYS_BELOW_7D_MEAN', 'LUFTTEMPERATUR_NO_DAYS_BELOW_7D_MEAN']

In [491...]: df.head()

Out[491]:	Montag	Dienstag	Mittwoch	Donnerstag	Freitag	Samstag	Sonntag	MONTH_JAN	MONTH_FEB	MONTH_MAR	MONTH_APR	MONTH_MAY
0	0	0	0	0	1	0	0	0	0	0	0	0
1	0	0	0	0	0	1	0	0	0	0	0	0
2	0	0	0	0	0	0	1	0	0	0	0	0
3	1	0	0	0	0	0	0	0	0	0	0	0
4	0	1	0	0	0	0	0	0	0	0	0	0

Poisson Regression

For the first iteration, we fit a Poisson Regression model to each menu item to predict sales, using the same features for each model.

In [408...]

```
from sklearn.model_selection import TimeSeriesSplit
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import PoissonRegressor

menu_items = ['CALAMARI', 'FISCH', 'GARNELEN', 'HAEHNCHEN', 'KOEFT', 'LAMM', 'STEAK']

X = df.drop(columns=menu_items)
features = X
y = df[menu_items]

# Scale features
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Set up TimeSeriesSplit
tscv = TimeSeriesSplit(n_splits=5)

for item in menu_items:
    train_scores = []
    test_scores = []

    for train_index, test_index in tscv.split(X):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y.iloc[train_index][item], y.iloc[test_index][item]
```

```
model = PoissonRegressor(max_iter=200)
model.fit(X_train, y_train)

y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)

train_mse = mean_squared_error(y_train, y_train_pred)
test_mse = mean_squared_error(y_test, y_test_pred)
test_scores.append(test_mse)
train_scores.append(train_mse)

avg_train_mse = sum(train_scores) / len(train_scores)
avg_test_mse = sum(test_scores) / len(test_scores)

train_mse = mean_squared_error(y_train, y_train_pred)
test_mse = mean_squared_error(y_test, y_test_pred)

print(f'{item} Avg Train MSE: {avg_train_mse:.4f}, Avg Test MSE: {avg_test_mse:.4f}\n')
```

CALAMARI Avg Train MSE: 4.3733, Avg Test MSE: 11.8854

FISCH Avg Train MSE: 4.3431, Avg Test MSE: 8.2688

GARNELEN Avg Train MSE: 9.3740, Avg Test MSE: 29.7110

HAEHNCHEN Avg Train MSE: 38.6646, Avg Test MSE: 287.8877

KOEFTE Avg Train MSE: 26.5360, Avg Test MSE: 140.8644

LAMM Avg Train MSE: 45.5634, Avg Test MSE: 217.1821

STEAK Avg Train MSE: 32.8859, Avg Test MSE: 357.1320

Results analysis: MSE: We get high values of mean squared error for haehnchen (chicken), lamm (lamb), steak. We also see a large overfitting as the test MSE for all categories are significantly higher than the train MSE. I would like to further look into why these values are high. This could possibly be caused by these items having more variabilities in sales or the features not capturing the nuances of these items as well

The model performs relatively well in predicting the sales of calamari, fisch (fish), and garnelen (shrimp), outputting a lower MSE

Train vs test scores: We can see for some of the menu items, like fisch (fish), garnelen (shrimp), haehnchen(chicken) that they perform better for the training set compared to the test set.

Hyperparameter Tuning

To help the model with overfitting, we can add a regularization method. The PoissonRegressor class allows for L2 regularization. We can also adjust the max_iter value to tune convergence.

GridSearchCV is a cross-validation technique that we will use to find the optimal parameter values in a grid. We will use this library to tune our hyperparameters, finding the optimal values.

In [409...]

```
param_grid = {
    'alpha': [0.01, 0.1, 1.0, 10.0, 100.0],
    'max_iter': [1000, 1500, 2000, 2500, 3000],
}

for item in menu_items:
    y_item = y[item]

    model = PoissonRegressor()
    grid_search = GridSearchCV(
        estimator=model,
        param_grid=param_grid,
        cv=tscv,
        scoring='neg_mean_squared_error',
        n_jobs=-1,
    )

    grid_search.fit(X, y_item)

    # Get best parameters and model
    best_params = grid_search.best_params_
    print(f"Best parameters for {item} :{best_params}")
    best_model = grid_search.best_estimator_

    split_index = int(len(df) * 0.8)
    X_train = X[:split_index]
    X_test = X[split_index:]
    y_train = y_item[:split_index]
    y_test = y_item[split_index:]
```

```
y_train_pred = best_model.predict(X_train)
y_pred = best_model.predict(X_test)

train_mse = mean_squared_error(y_train, y_train_pred)
test_mse = mean_squared_error(y_test, y_pred)

print(f"Train MSE: {train_mse:.4f}")
print(f"Test MSE: {test_mse:.4f}\n")
```

Best parameters for CALAMARI :{'alpha': 100.0, 'max_iter': 1000}

Train MSE: 8.1955

Test MSE: 4.5404

Best parameters for FISCH :{'alpha': 10.0, 'max_iter': 1000}

Train MSE: 6.3633

Test MSE: 4.2153

Best parameters for GARNELEN :{'alpha': 10.0, 'max_iter': 1000}

Train MSE: 14.8100

Test MSE: 14.3984

Best parameters for HAEHNCHEN :{'alpha': 100.0, 'max_iter': 1000}

Train MSE: 86.5244

Test MSE: 109.7315

Best parameters for KOEFTTE :{'alpha': 100.0, 'max_iter': 1000}

Train MSE: 57.1914

Test MSE: 68.5894

Best parameters for LAMM :{'alpha': 100.0, 'max_iter': 1000}

Train MSE: 102.5174

Test MSE: 91.9662

Best parameters for STEAK :{'alpha': 100.0, 'max_iter': 1000}

Train MSE: 67.4207

Test MSE: 56.7634

The MSE lowered significantly for all the categories. However, the model still seems to be overfitting in some categories.

Feature Evaluation

We will perform feature selection to try to mitigate the overfitting. We can use the coefficients to analyze which features have the strongest importance in the model's performance, and select the top 10% of them. This will reduce noise and potentially help

with overfitting.

```
In [52]: coefficients = best_model.coef_
feature_importance = abs(coefficients)

importance_df = pd.DataFrame({
    'Feature': features.columns,
    'Coefficient': coefficients,
    'Importance': feature_importance
}).sort_values(by='Importance', ascending=False)

top_features = importance_df.head(int(273* 0.10))['Feature'].tolist()
print(top_features)

['Samstag', 'STEAK_DEMAND_T7', 'HAEHNCHEN_DEMAND_T7', 'Sonntag', 'CALAMARI_MEAN_SAME_WDAY_DEMANDS_W2', 'KOEFT_E_DEMAND_T7', 'LAMM_DEMAND_T7', 'YEAR_2013', 'CALAMARI_DEMAND_T7', 'WEEKEND', 'FISCH_MEAN_SAME_WDAY_DEMANDS_W2', 'GARNELEN_DEMAND_T7', 'FISCH_DEMAND_T7', 'Montag', 'LUFTTEMPERATUR', 'MONTH_SEP', 'LAMM_DEMAND_T5', 'MONTH_NOV', 'Dienstag', 'ISHOLIDAY', 'YEAR_2015', 'LAMM_DEMAND_T2', 'STEAK_DEMAND_T1', 'SONNE_HML_OVER_T7', 'LAMM_DEMAND_T3', 'STEAK_DEMAND_T6', 'CALAMARI_DEMAND_T1']
```

```
In [53]: top_feature_indices = [features.columns.get_loc(f) for f in top_features]
X_top_features = X[:, top_feature_indices]

for item in menu_items:
    y_item = y[item]

    model = PoissonRegressor()
    grid_search = GridSearchCV(
        estimator=model,
        param_grid=param_grid,
        cv=tscv, # Use TimeSeriesSplit for cross-validation
        scoring='neg_mean_squared_error', # Metric to evaluate the model
        n_jobs=-1 # Use all available CPUs
    )

    grid_search.fit(X_top_features, y_item)

    # Get best parameters and model
    best_params = grid_search.best_params_
    print(f"Best parameters for {item} with top features: {best_params}")
    best_model = grid_search.best_estimator_

    # Evaluate the best model
    split_index = int(len(df) * 0.8)
    X_train = X_top_features[:split_index]
```

```
X_test = X_top_features[split_index:]
y_train = y_item[:split_index]
y_test = y_item[split_index:]

y_train_pred = best_model.predict(X_train)
y_pred = best_model.predict(X_test)

train_mse = mean_squared_error(y_train, y_train_pred)
test_mse = mean_squared_error(y_test, y_pred)

print(f"Train MSE with top features: {train_mse:.4f}")
print(f"Test MSE with top features: {test_mse:.4f}\n")
```

```
Best parameters for CALAMARI with top features: {'alpha': 10.0, 'max_iter': 1000}
Train MSE with top features: 7.2896
Test MSE with top features: 4.1257
```

```
Best parameters for FISCH with top features: {'alpha': 1.0, 'max_iter': 1000}
Train MSE with top features: 6.3992
Test MSE with top features: 4.3659
```

```
Best parameters for GARNELEN with top features: {'alpha': 10.0, 'max_iter': 1000}
Train MSE with top features: 16.6125
Test MSE with top features: 15.9592
```

```
Best parameters for HAEHNCHEN with top features: {'alpha': 10.0, 'max_iter': 1000}
Train MSE with top features: 70.6049
Test MSE with top features: 105.0817
```

```
Best parameters for KOEFTET with top features: {'alpha': 10.0, 'max_iter': 1000}
Train MSE with top features: 49.2375
Test MSE with top features: 68.9013
```

```
Best parameters for LAMM with top features: {'alpha': 10.0, 'max_iter': 1000}
Train MSE with top features: 87.6822
Test MSE with top features: 86.8723
```

```
Best parameters for STEAK with top features: {'alpha': 10.0, 'max_iter': 1000}
Train MSE with top features: 55.7043
Test MSE with top features: 50.4318
```

Diminishing the number of features to the top 10% most important features seemed to help a lot with overfitting. The MSE also went down a bit from the last model that we predicted from.

Random Forest

```
In [54]: # Define the parameter grid for hyperparameter tuning
params = {
    'n_estimators': [50, 100, 150],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt', 'log2']
}

tscv = TimeSeriesSplit(n_splits=5)

for item in menu_items:
    y = df[item]

    mse_scores_train = []
    mse_scores_test = []

    all_y_train_pred = []
    all_y_test_pred = []
    all_y_train_true = []
    all_y_test_true = []

    for train_index, test_index in tscv.split(X_top_features):
        X_train, X_test = X_top_features[train_index], X_top_features[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

        model = RandomForestRegressor(random_state=42)
        grid_search = GridSearchCV(
            estimator=model,
            param_grid=params,
            cv=3,
            scoring='neg_mean_squared_error',
            n_jobs=-1
        )
        grid_search.fit(X_train, y_train)

        best_model = grid_search.best_estimator_

        y_train_pred = best_model.predict(X_train)
        y_test_pred = best_model.predict(X_test)
```

```
        mse_train = mean_squared_error(y_train, y_train_pred)
        mse_test = mean_squared_error(y_test, y_test_pred)

        mse_scores_train.append(mse_train)
        mse_scores_test.append(mse_test)

        all_y_train_pred.extend(y_train_pred)
        all_y_test_pred.extend(y_test_pred)
        all_y_train_true.extend(y_train)
        all_y_test_true.extend(y_test)

        avg_mse_train = sum(mse_scores_train) / len(mse_scores_train)
        avg_mse_test = sum(mse_scores_test) / len(mse_scores_test)

        print(f"Best parameters for {item} :{grid_search.best_params_}")
        print(f"Average Train MSE: {avg_mse_train:.4f}")
        print(f"Average Test MSE: {avg_mse_test:.4f}\n")
```

```
Best parameters for CALAMARI :{'max_depth': None, 'max_features': 'log2', 'min_samples_leaf': 4, 'min_samples_split': 2, 'n_estimators': 150}
Average Train MSE: 4.4584
Average Test MSE: 6.0254

Best parameters for FISCH :{'max_depth': 10, 'max_features': 'log2', 'min_samples_leaf': 4, 'min_samples_split': 2, 'n_estimators': 50}
Average Train MSE: 3.0892
Average Test MSE: 6.3536

Best parameters for GARNELEN :{'max_depth': 10, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 10, 'n_estimators': 150}
Average Train MSE: 7.0995
Average Test MSE: 18.5810

Best parameters for HAEHNCHEN :{'max_depth': 20, 'max_features': 'log2', 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}
Average Train MSE: 19.8043
Average Test MSE: 94.8375

Best parameters for KOEFTETE :{'max_depth': 10, 'max_features': 'log2', 'min_samples_leaf': 1, 'min_samples_split': 10, 'n_estimators': 150}
Average Train MSE: 20.6918
Average Test MSE: 56.1709

Best parameters for LAMM :{'max_depth': None, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 10, 'n_estimators': 100}
Average Train MSE: 26.8727
Average Test MSE: 104.6050

Best parameters for STEAK :{'max_depth': 10, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 150}
Average Train MSE: 21.4422
Average Test MSE: 66.1955
```

Random Forest Regression gives us similar results as Poisson Regression, but we can see here that overfitting is a much larger problem. This is a common issue because of Random Forest's flexibility, which allows the model to capture complex patterns but also makes it more susceptible to capturing noise within the data, causing it to overfit to the training data.

XGBoost

In [55]:

```
from sklearn.model_selection import TimeSeriesSplit, GridSearchCV
from sklearn.metrics import mean_squared_error
from xgboost import XGBRegressor

param_grid = {
    'n_estimators': [100, 200],
    'learning_rate': [0.01, 0.1],
    'max_depth': [3, 5],
    'subsample': [0.7, 0.8],
    'colsample_bytree': [0.7, 0.8],
    'min_child_weight': [1, 3],
    'gamma': [0, 0.1],
    'lambda': [0, 1],
    'alpha': [0, 1]
}

for item in menu_items:
    y = df[item]

    tscv = TimeSeriesSplit(n_splits=5)

    model = XGBRegressor(objective='reg:squarederror')
    grid_search = GridSearchCV(
        estimator=model,
        param_grid=param_grid,
        cv=tscv,
        scoring='neg_mean_squared_error',
        n_jobs=-1,
    )

    grid_search.fit(X_top_features, y)

    best_params = grid_search.best_params_
    print(f"Best parameters for {item}: {best_params}")
    best_model = grid_search.best_estimator_

    y_pred = best_model.predict(X_top_features[tscv.get_n_splits():])

    y_train = y[:-len(y_pred)]
    y_test = y[-len(y_pred):]

    train_mse = mean_squared_error(y_train, best_model.predict(X_top_features[:-len(y_pred)]))
    test_mse = mean_squared_error(y_test, y_pred)
```

```
print(f"Train MSE: {train_mse}")
print(f"Test MSE: {test_mse}\n")
```

```
Best parameters for CALAMARI: {'alpha': 0, 'colsample_bytree': 0.7, 'gamma': 0, 'lambda': 1, 'learning_rate': 0.01, 'max_depth': 3, 'min_child_weight': 1, 'n_estimators': 200, 'subsample': 0.7}
```

```
Train MSE: 2.9590394983714416
```

```
Test MSE: 5.164862514571149
```

```
Best parameters for FISCH: {'alpha': 1, 'colsample_bytree': 0.8, 'gamma': 0.1, 'lambda': 1, 'learning_rate': 0.01, 'max_depth': 3, 'min_child_weight': 3, 'n_estimators': 200, 'subsample': 0.7}
```

```
Train MSE: 0.8395606756522739
```

```
Test MSE: 5.288939551645853
```

```
Best parameters for GARNELEN: {'alpha': 0, 'colsample_bytree': 0.8, 'gamma': 0, 'lambda': 0, 'learning_rate': 0.01, 'max_depth': 3, 'min_child_weight': 3, 'n_estimators': 100, 'subsample': 0.7}
```

```
Train MSE: 24.848164786230154
```

```
Test MSE: 16.154454542670873
```

```
Best parameters for HAEHNCHEN: {'alpha': 1, 'colsample_bytree': 0.7, 'gamma': 0, 'lambda': 1, 'learning_rate': 0.01, 'max_depth': 3, 'min_child_weight': 1, 'n_estimators': 200, 'subsample': 0.8}
```

```
Train MSE: 31.219736649170226
```

```
Test MSE: 67.02939114746582
```

```
Best parameters for KOEFTE: {'alpha': 0, 'colsample_bytree': 0.7, 'gamma': 0.1, 'lambda': 1, 'learning_rate': 0.01, 'max_depth': 5, 'min_child_weight': 1, 'n_estimators': 200, 'subsample': 0.7}
```

```
Train MSE: 15.787731704912584
```

```
Test MSE: 32.66704601223124
```

```
Best parameters for LAMM: {'alpha': 1, 'colsample_bytree': 0.8, 'gamma': 0.1, 'lambda': 0, 'learning_rate': 0.01, 'max_depth': 3, 'min_child_weight': 3, 'n_estimators': 200, 'subsample': 0.8}
```

```
Train MSE: 67.75551329617738
```

```
Test MSE: 76.06063484390435
```

```
Best parameters for STEAK: {'alpha': 1, 'colsample_bytree': 0.8, 'gamma': 0.1, 'lambda': 1, 'learning_rate': 0.01, 'max_depth': 5, 'min_child_weight': 1, 'n_estimators': 200, 'subsample': 0.8}
```

```
Train MSE: 12.026724873802596
```

```
Test MSE: 30.488669640088954
```

So far, XGBoost has given the best results (lowest Test MSE), though it does overfit on the data.

Time Series Forecasting

We will use the SARIMA (Seasonal autoregressive integrated moving average) model to predict future sales. This model is designed to capture trends and seasonality. Naturally, restaurant sales follow a period of seasonality trends with slow and busy seasons, correlating with the time of year (holidays), and weekdays vs. weekends.

SARIMA(p,d,q)×(P,D,Q)m

Non-Seasonal Component (p, d, q):

p: Autoregressive (AR) part p is the number of lag observations in the model (number of autoregressive terms).

d: Differencing (I) part d is the number of times the raw observations are differenced to make the time series stationary.

q Moving Average (MA) part: The parameter q is the size of the moving average window (number of lagged forecast errors).

Seasonal Component (P, D, Q):

P: Seasonal Autoregressive (SAR) part: The parameter P is the number of seasonal autoregressive terms.

D: Seasonal Differencing (SI) part: The parameter D is the number of seasonal differences.

Q: Seasonal Moving Average (SMA) part: The parameter Q is the number of seasonal moving average terms.

m: Seasonal Period (m): The number of periods per season (e.g., 7 for weekly trends).

```
In [504...]: # Convert DEMAND_DATE into datetime
df_time_series = pd.concat([df_time_series, df], axis=1)
df_time_series['DEMAND_DATE'] = pd.to_datetime(df_time_series['DEMAND_DATE'], format='%m.%d.%Y')
df_time_series = df_time_series.set_index('DEMAND_DATE')
df_time_series.head()

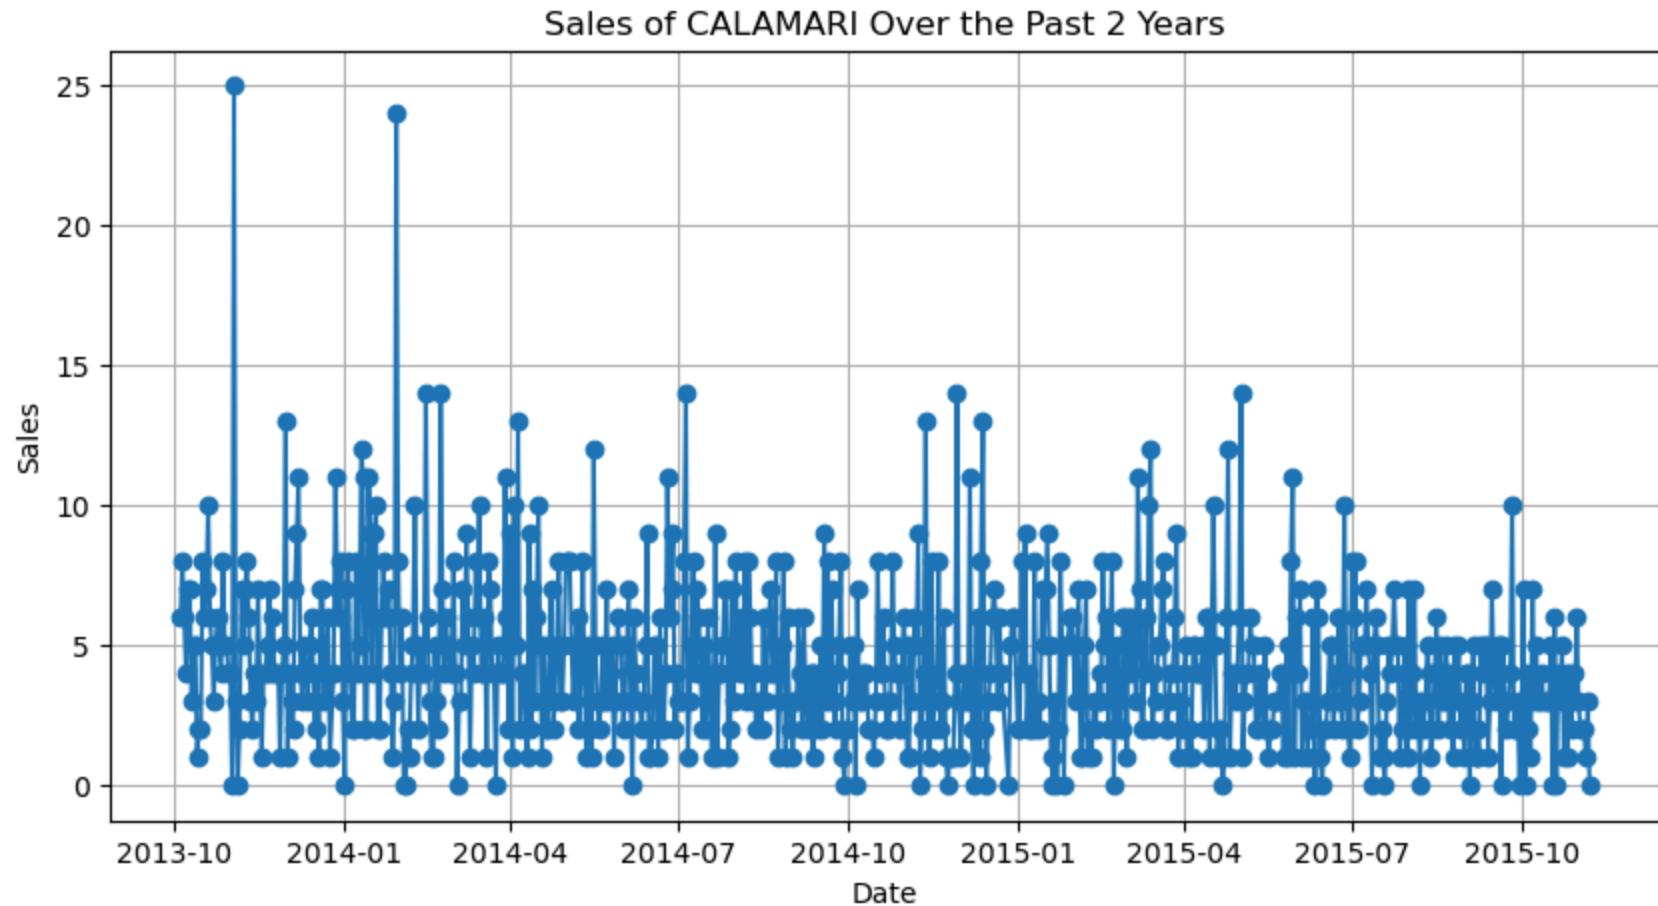
df_exog = df_time_series[['LUFTTEMPERATUR', 'ISHOLIDAY', 'WEEKEND']]
df_time_series = df_time_series[menu_items]
```

```
In [505...]: plt.figure(figsize=(14,5))

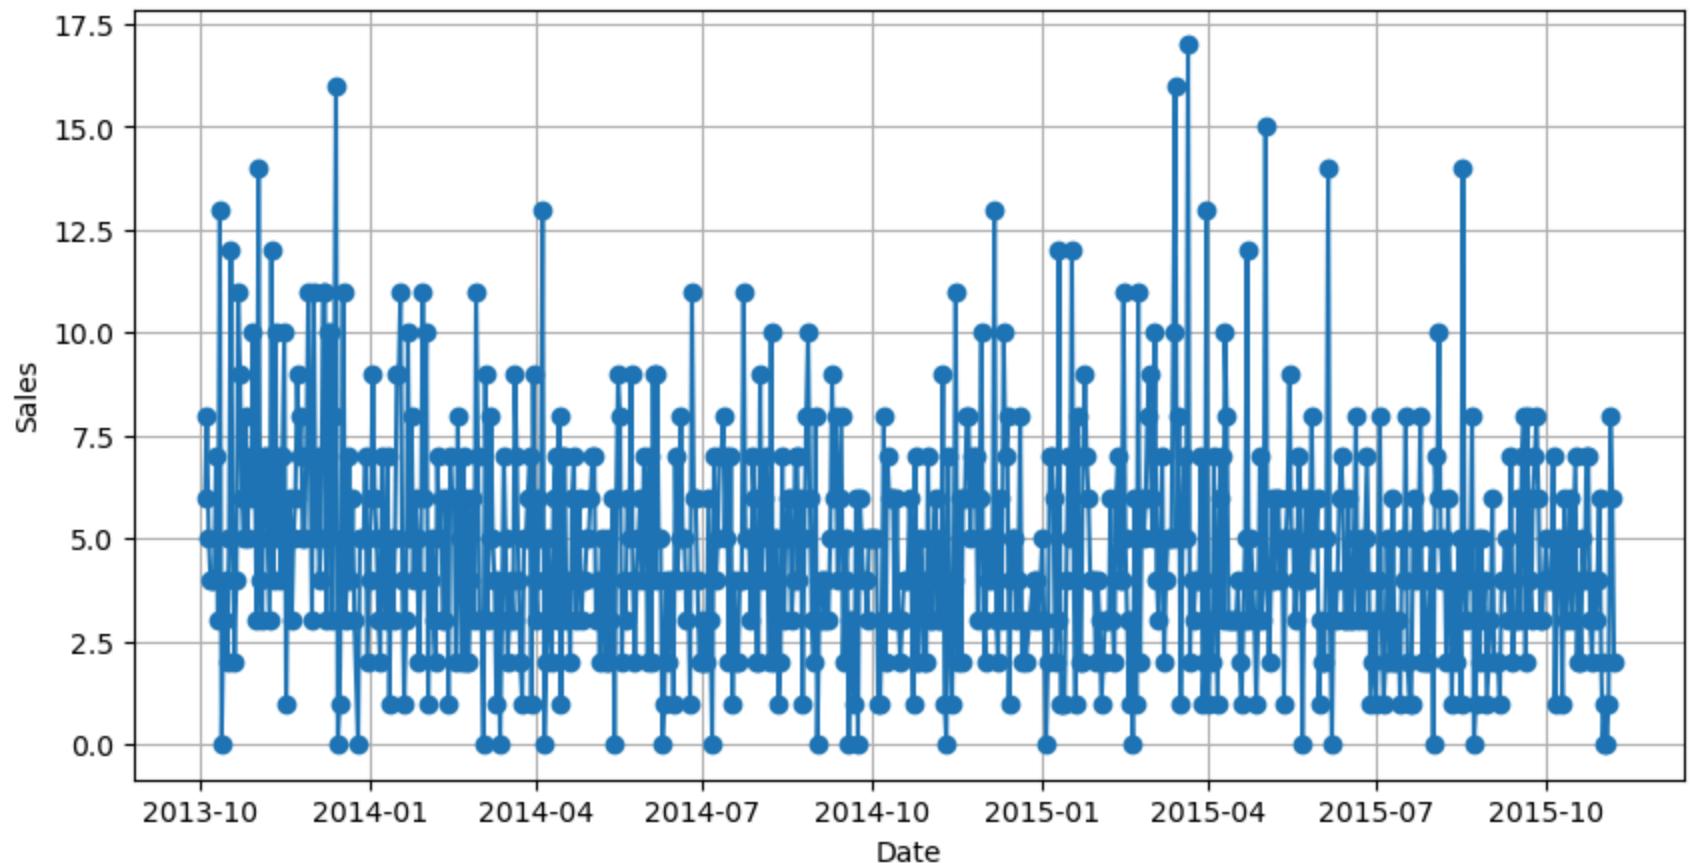
for item in menu_items:
    plt.figure(figsize=(10, 5))
    plt.plot(df_time_series.index, df_time_series[item], marker='o')
```

```
plt.title(f'Sales of {item} Over the Past 2 Years')
plt.xlabel('Date')
plt.ylabel('Sales')
plt.grid(True)
plt.show()
```

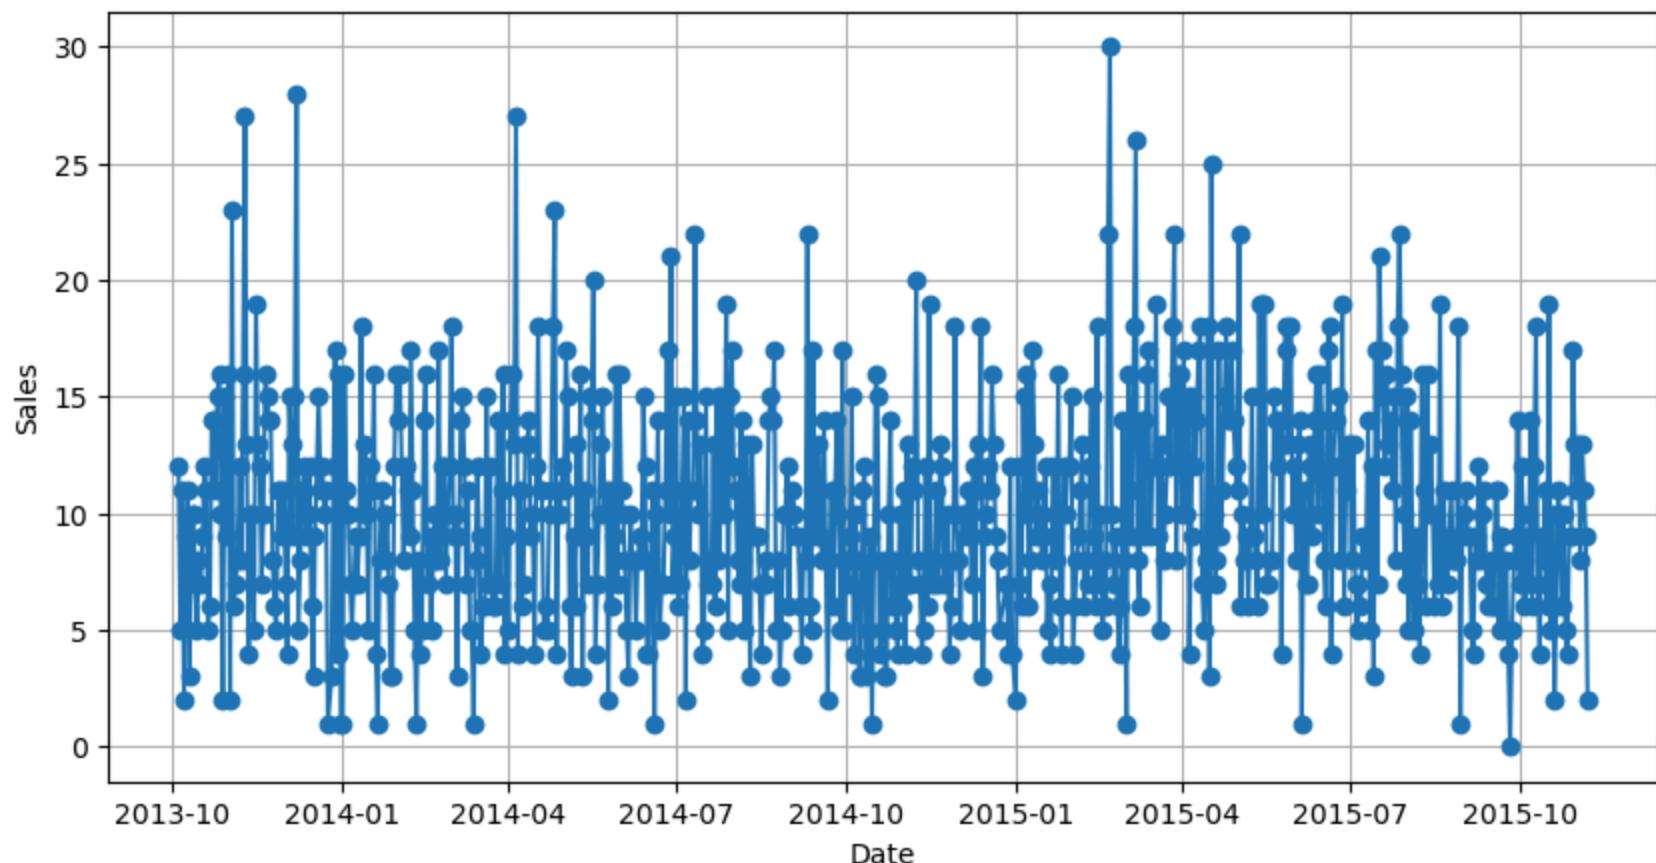
<Figure size 1400x500 with 0 Axes>



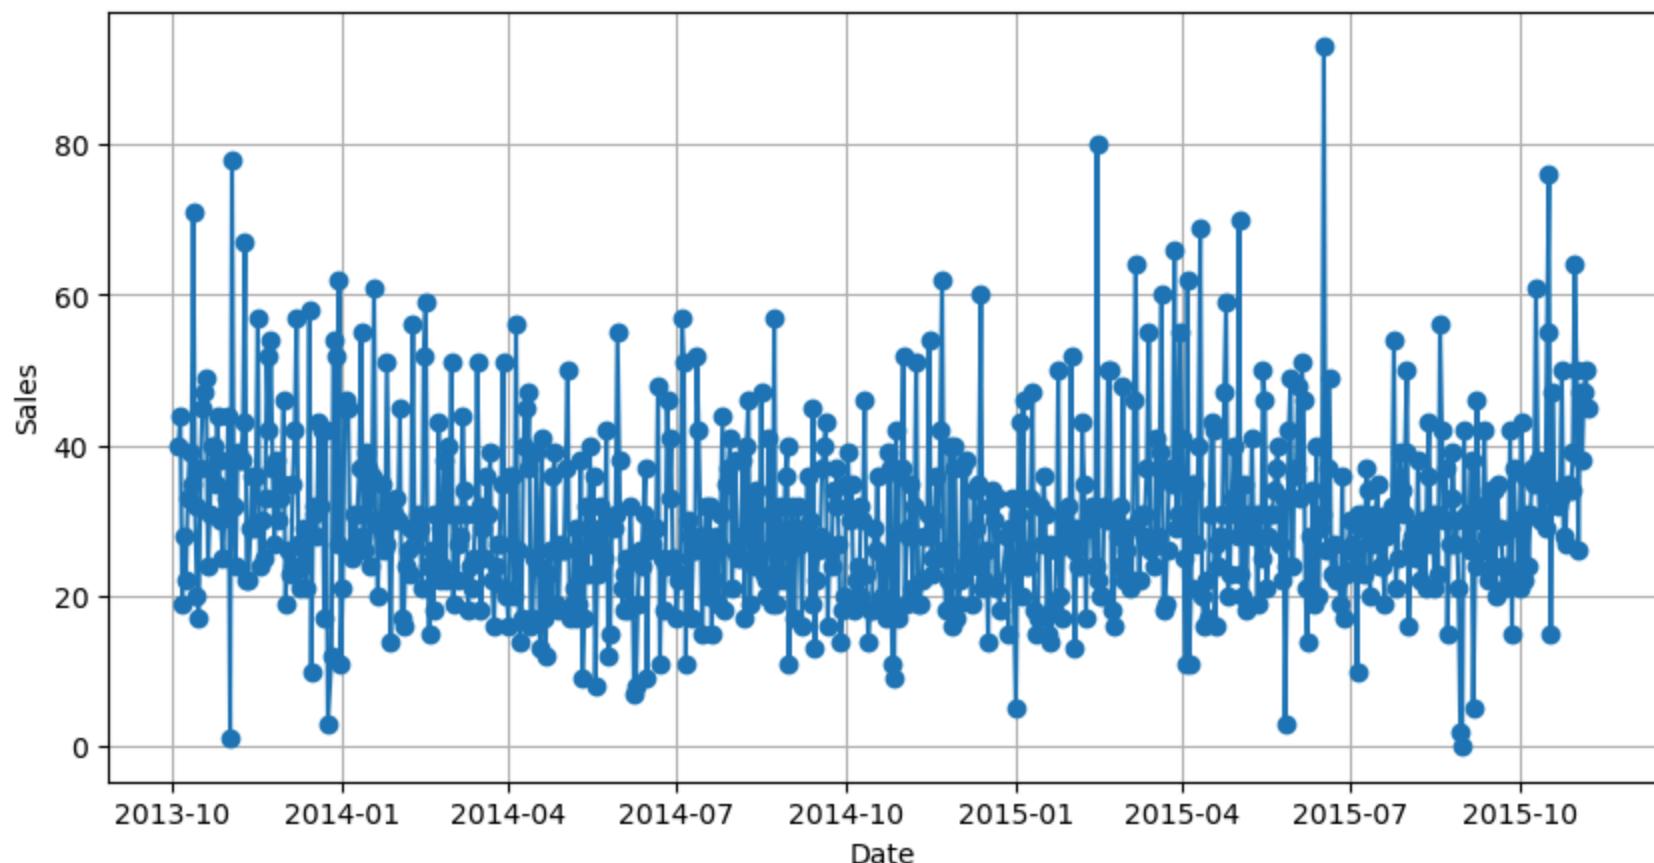
Sales of FISCH Over the Past 2 Years



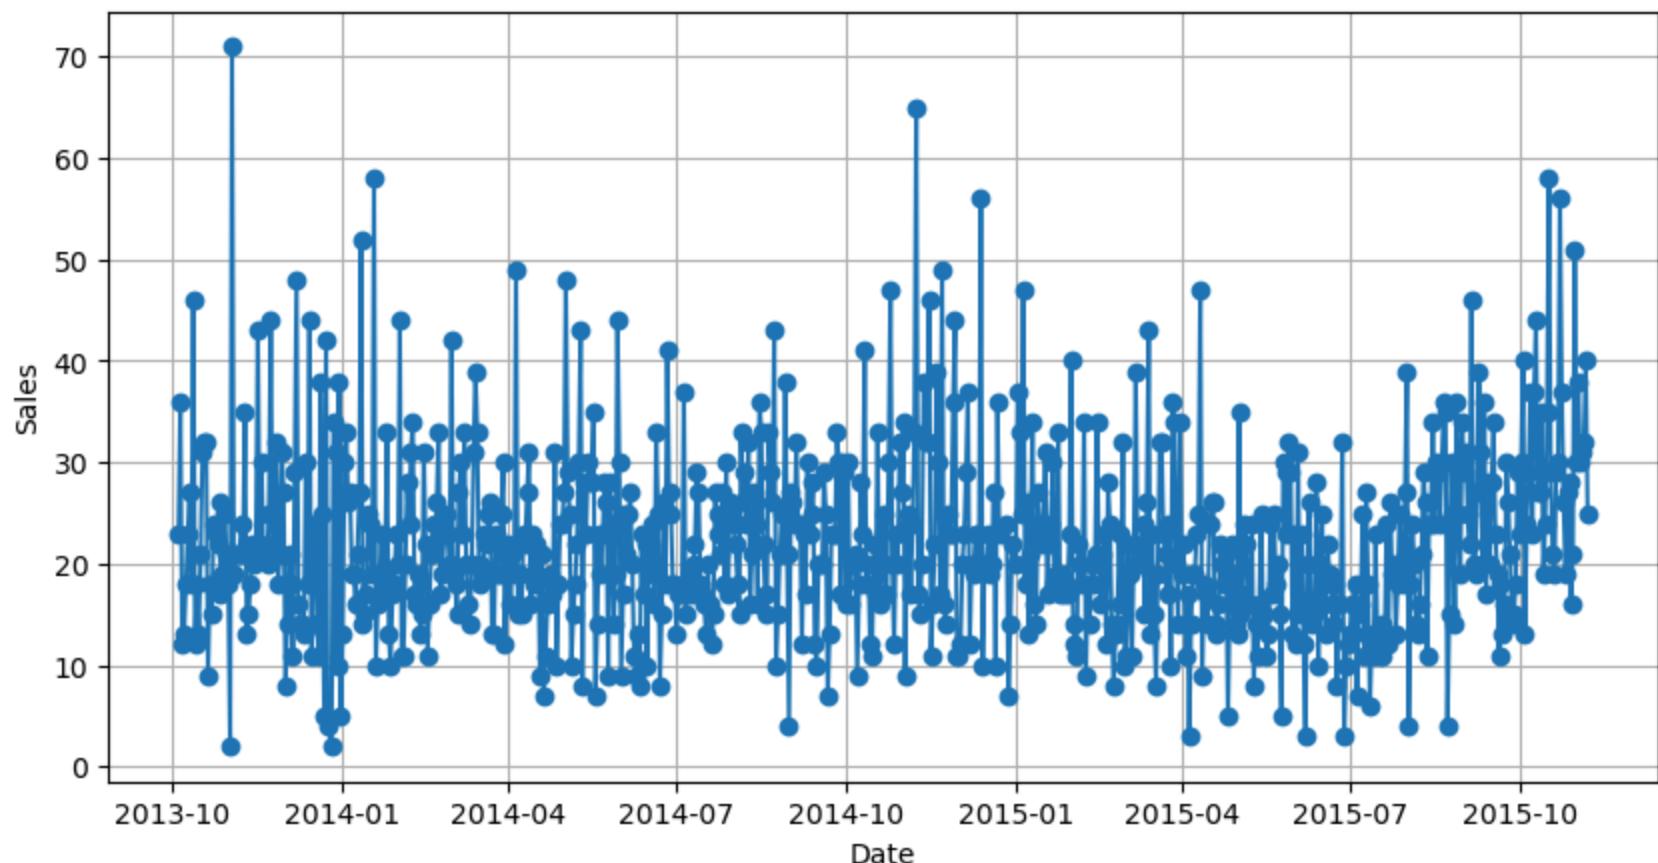
Sales of GARNELEN Over the Past 2 Years



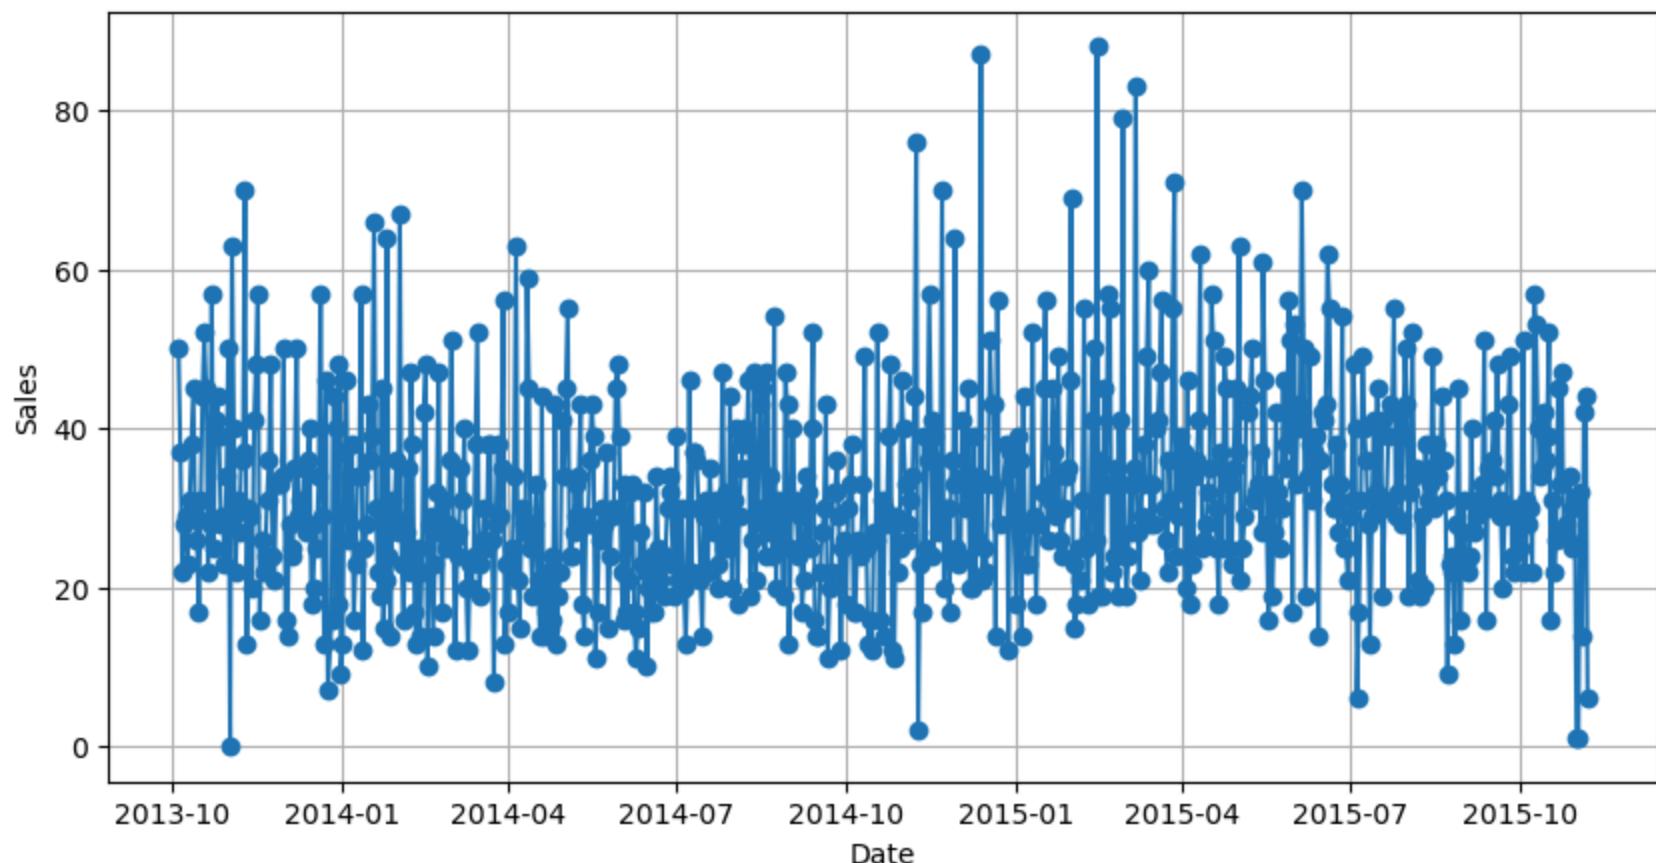
Sales of HAEHNCHEN Over the Past 2 Years



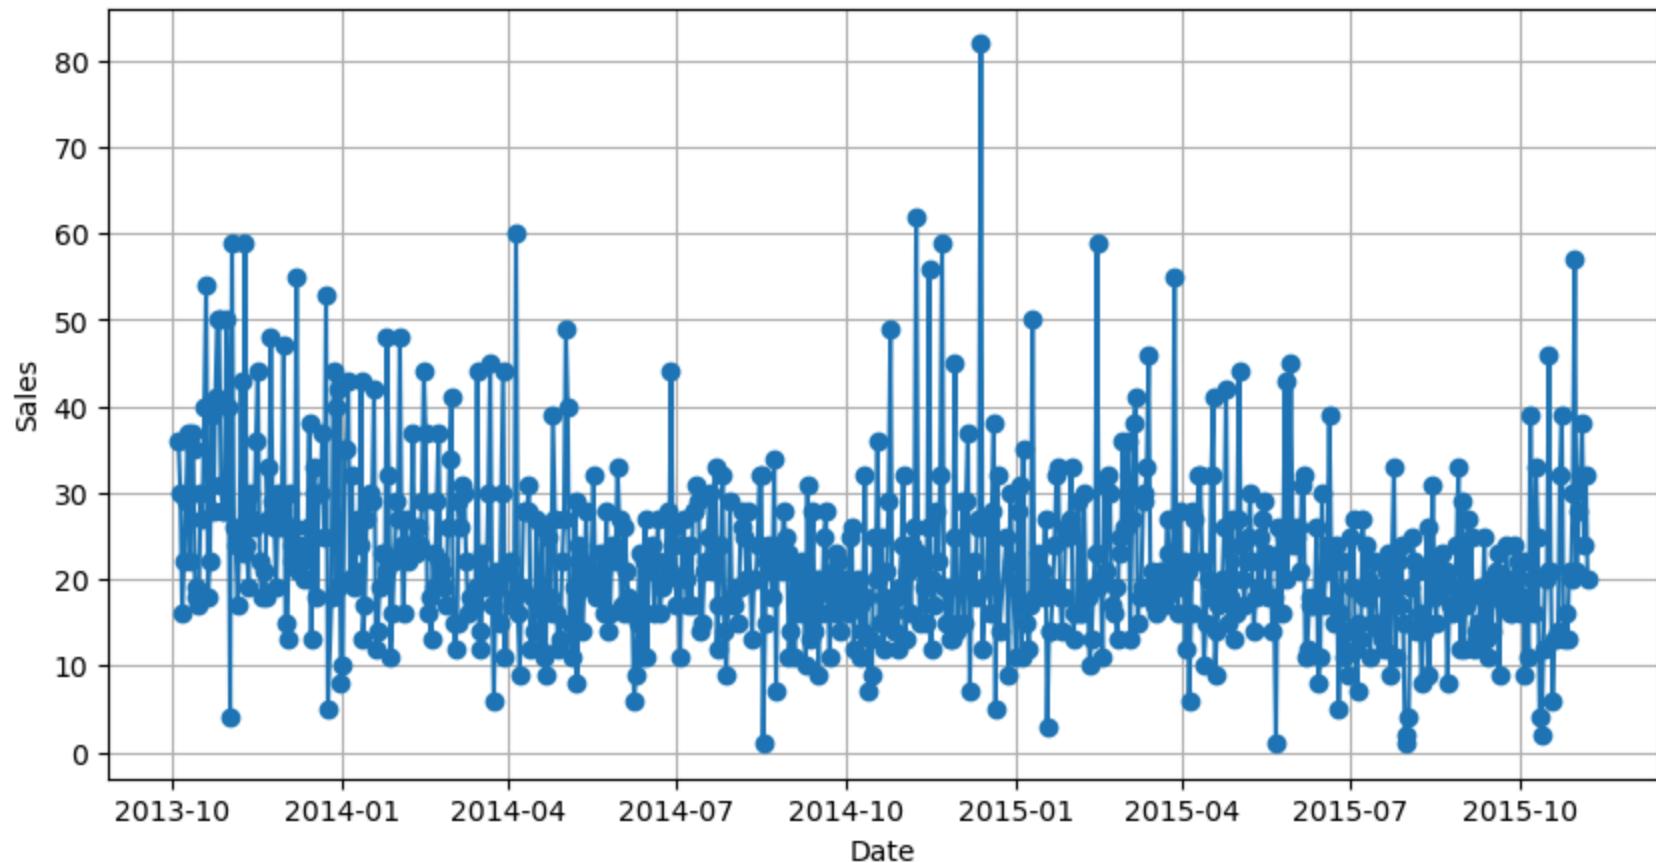
Sales of KOEFT Over the Past 2 Years



Sales of LAMM Over the Past 2 Years



Sales of STEAK Over the Past 2 Years



SARIMA model requires data to be stationary. We can use the Augmented Dickey-Fuller (ADF) Test to check if the data is stationary. (stationary if p-value < 0.05). A stationary time series means that the data will have a steady statistical trend across time.

Check Stationarity

```
In [417...]:  
def adf_test(series, signif=0.05):  
    result = adfuller(series)  
    return result[1] < signif  
  
for item in menu_items:  
    print(f"{item} stationary:", adf_test(df_time_series[item]))
```

```
CALAMARI stationary: True
FISCH stationary: True
GARNELEN stationary: True
HAEHNCHEN stationary: False
KOEFTET stationary: False
LAMM stationary: True
STEAK stationary: True
```

Haehnchen and Koefte are not stationary, so we need to difference the data.

```
In [420...]: not_stationary = ['HAEHNCHEN', 'KOEFTET']
for item in not_stationary:
    df_time_series[item] = df_time_series[item].diff().dropna()
    print(f"{item} stationary:", adf_test(df_diff[item]))
```

```
HAEHNCHEN stationary: True
KOEFTET stationary: True
```

Now all the menu items are stationary, so we can continue on to fit the model

Finding best parameters (p, d, q) and (P, D, Q)

To find the optimal values for p and q, we iteratively find the AIC values of a SARIMA model with p and q values within the range of 0 and 2, and pick the values that return the lowest model AIC value. The AIC value is a statistical relative measure of the quality of a model. A lower value indicates a better model fit.

```
In [423...]: import itertools
from statsmodels.tsa.statespace.sarimax import SARIMAX
import matplotlib.pyplot as plt

p_values = range(0, 2)
q_values = range(0, 2)
P_values = range(0, 2)
Q_values = range(0, 2)
seasonal_period = 7 # Weekly seasonality

def find_best_params(plot_diagnostics=True):
    best_params = {}

    for item in menu_items:
        print(f"Evaluating {item}")
        best_aic = float('inf')
```

```
best_p = best_q = best_P = best_Q = None

d = 1 if item in not_stationary else 0
D = 1 if item in not_stationary else 0

for p, q, P, Q in itertools.product(p_values, q_values, P_values, Q_values):
    try:
        model = SARIMAX(
            df[item],
            order=(p, d, q),
            seasonal_order=(P, D, Q, seasonal_period),
            enforce_stationarity=False,
            enforce_invertibility=False
        )
        sarima_model = model.fit(disp=False)

        if sarima_model.aic < best_aic:
            best_aic = sarima_model.aic
            best_p, best_q, best_P, best_Q = p, q, P, Q

    except Exception as e:
        print(f"Error with p={p}, d={d}, q={q}, P={P}, D={D}, Q={Q}: {e}")
        continue

    best_params[item] = {
        'p': best_p,
        'd': d,
        'q': best_q,
        'P': best_P,
        'D': D,
        'Q': best_Q
    }

    best_model = SARIMAX(
        df[item],
        order=(best_p, d, best_q),
        seasonal_order=(best_P, D, best_Q, seasonal_period),
        enforce_stationarity=False,
        enforce_invertibility=False
    )
    best_sarima_model = best_model.fit(disp=False)
    print(f"Best model for {item}: p={best_p}, d={d}, q={best_q}, P={best_P}, D={D}, Q={best_Q}")

    if plot_diagnostics:
        best_sarima_model.plot_diagnostics(figsize=(14, 8))
```

```

    plt.show()

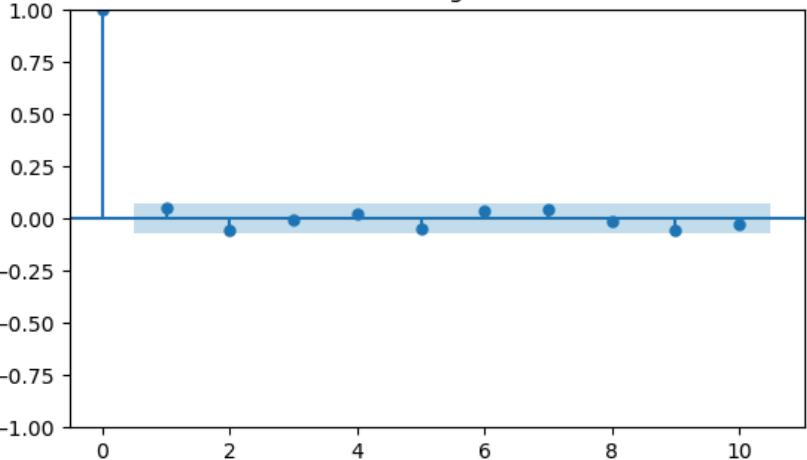
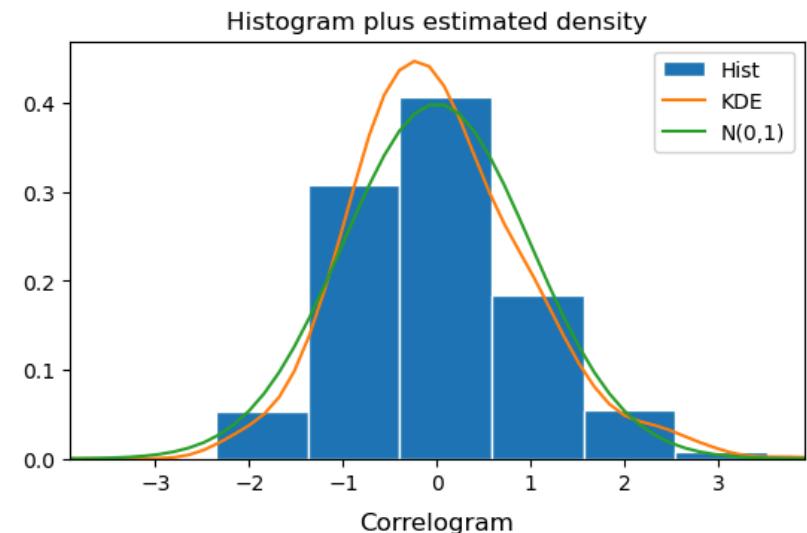
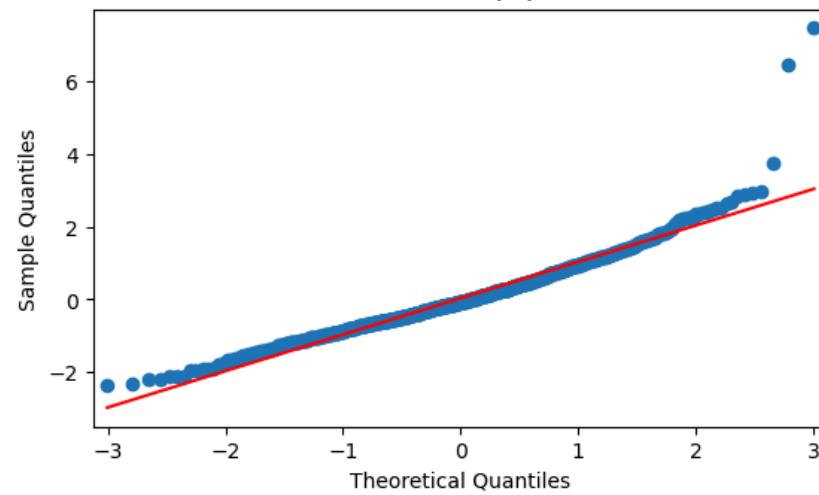
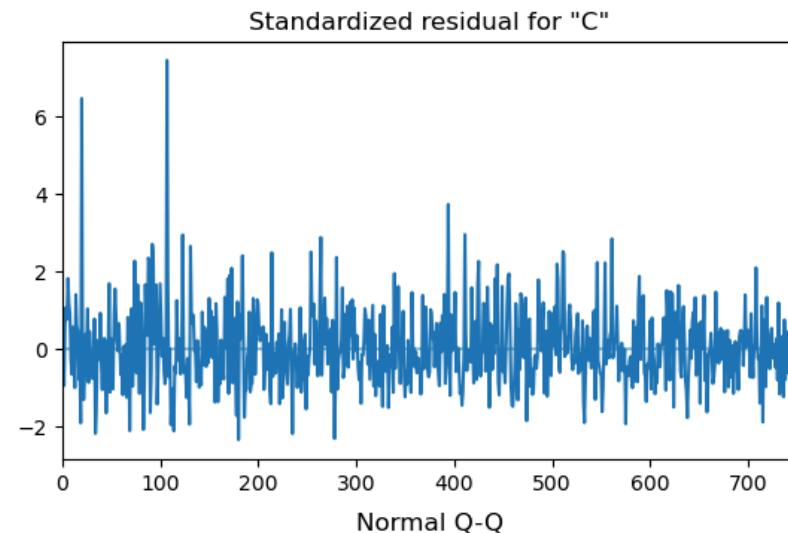
    return best_params

best_params = find_best_params()
print(best_params)

```

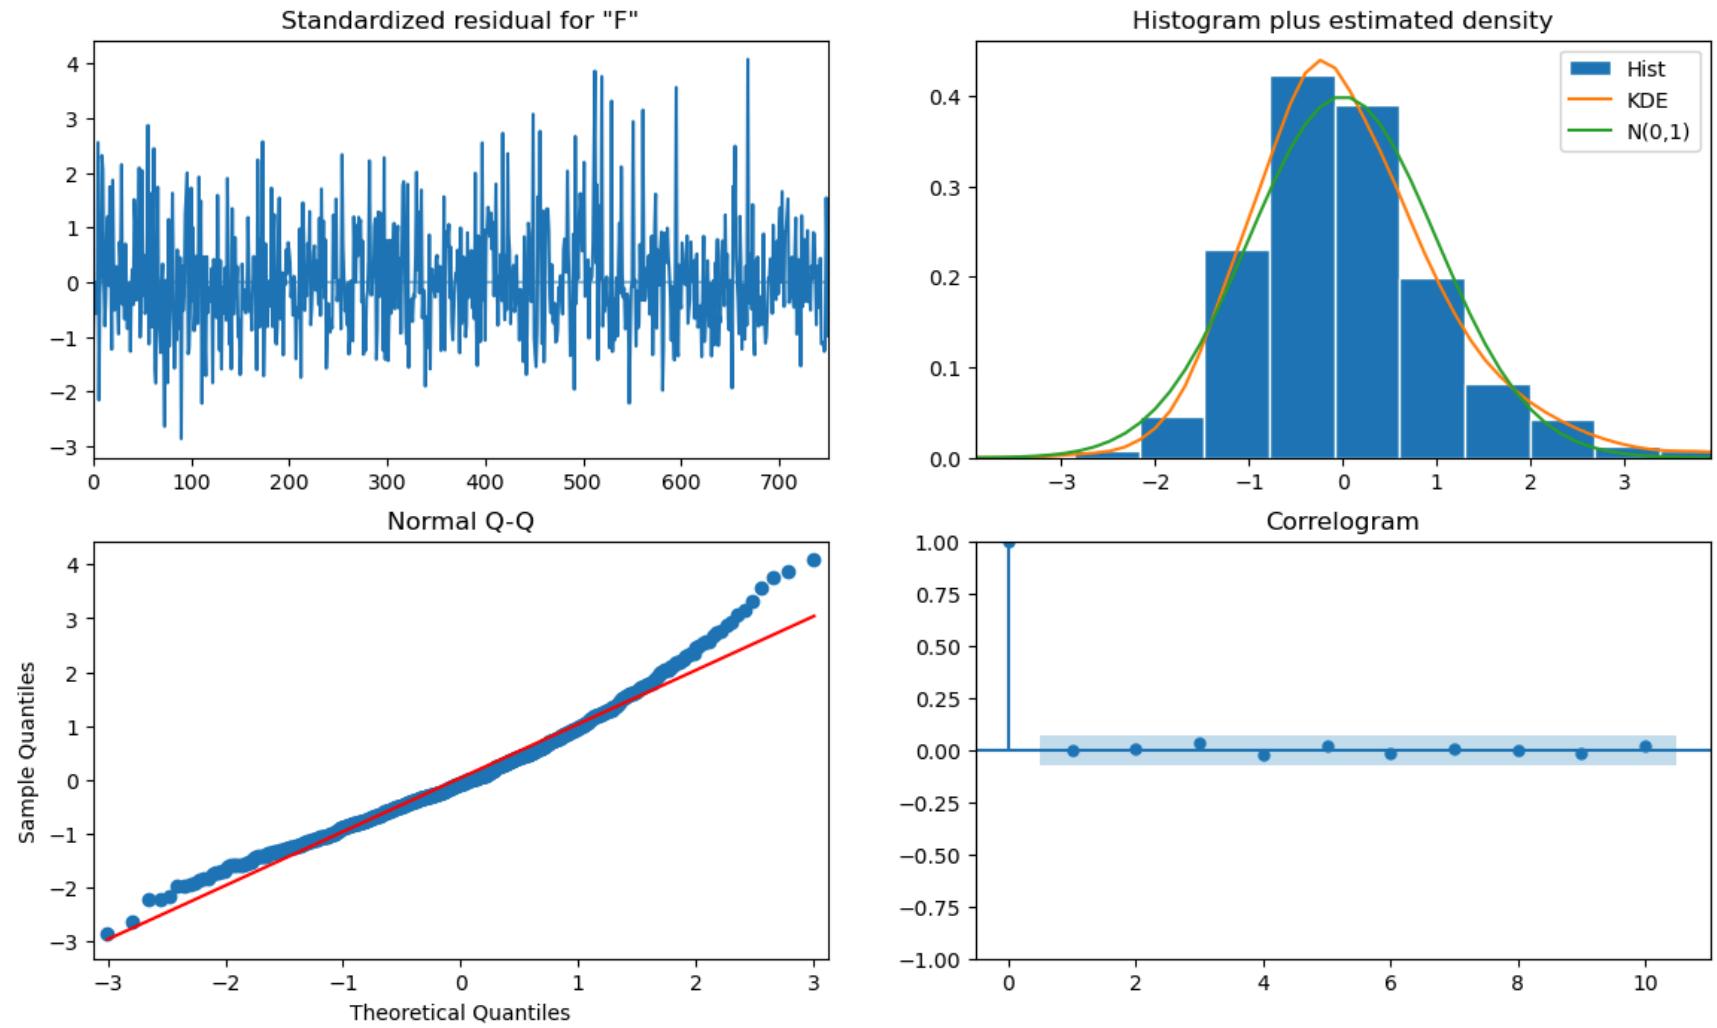
Evaluating CALAMARI

Best model for CALAMARI: p=1, d=0, q=1, P=1, D=0, Q=1



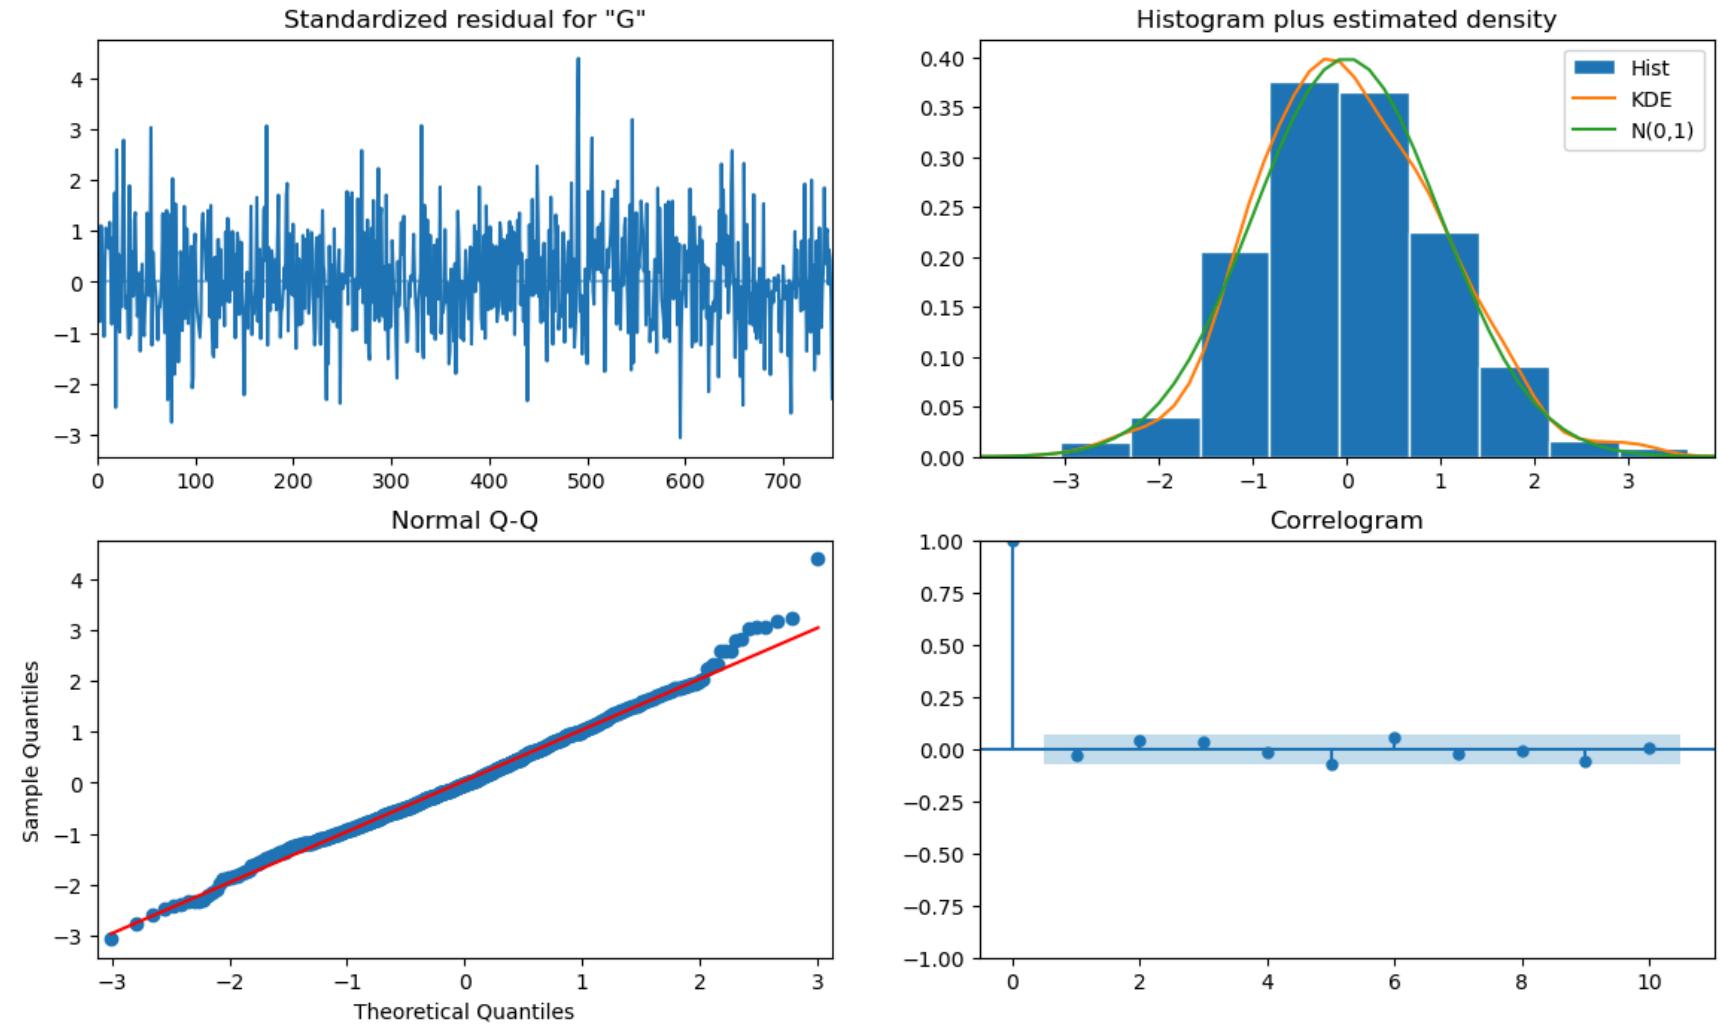
Evaluating FISCH

Best model for FISCH: p=0, d=0, q=1, P=1, D=0, Q=1



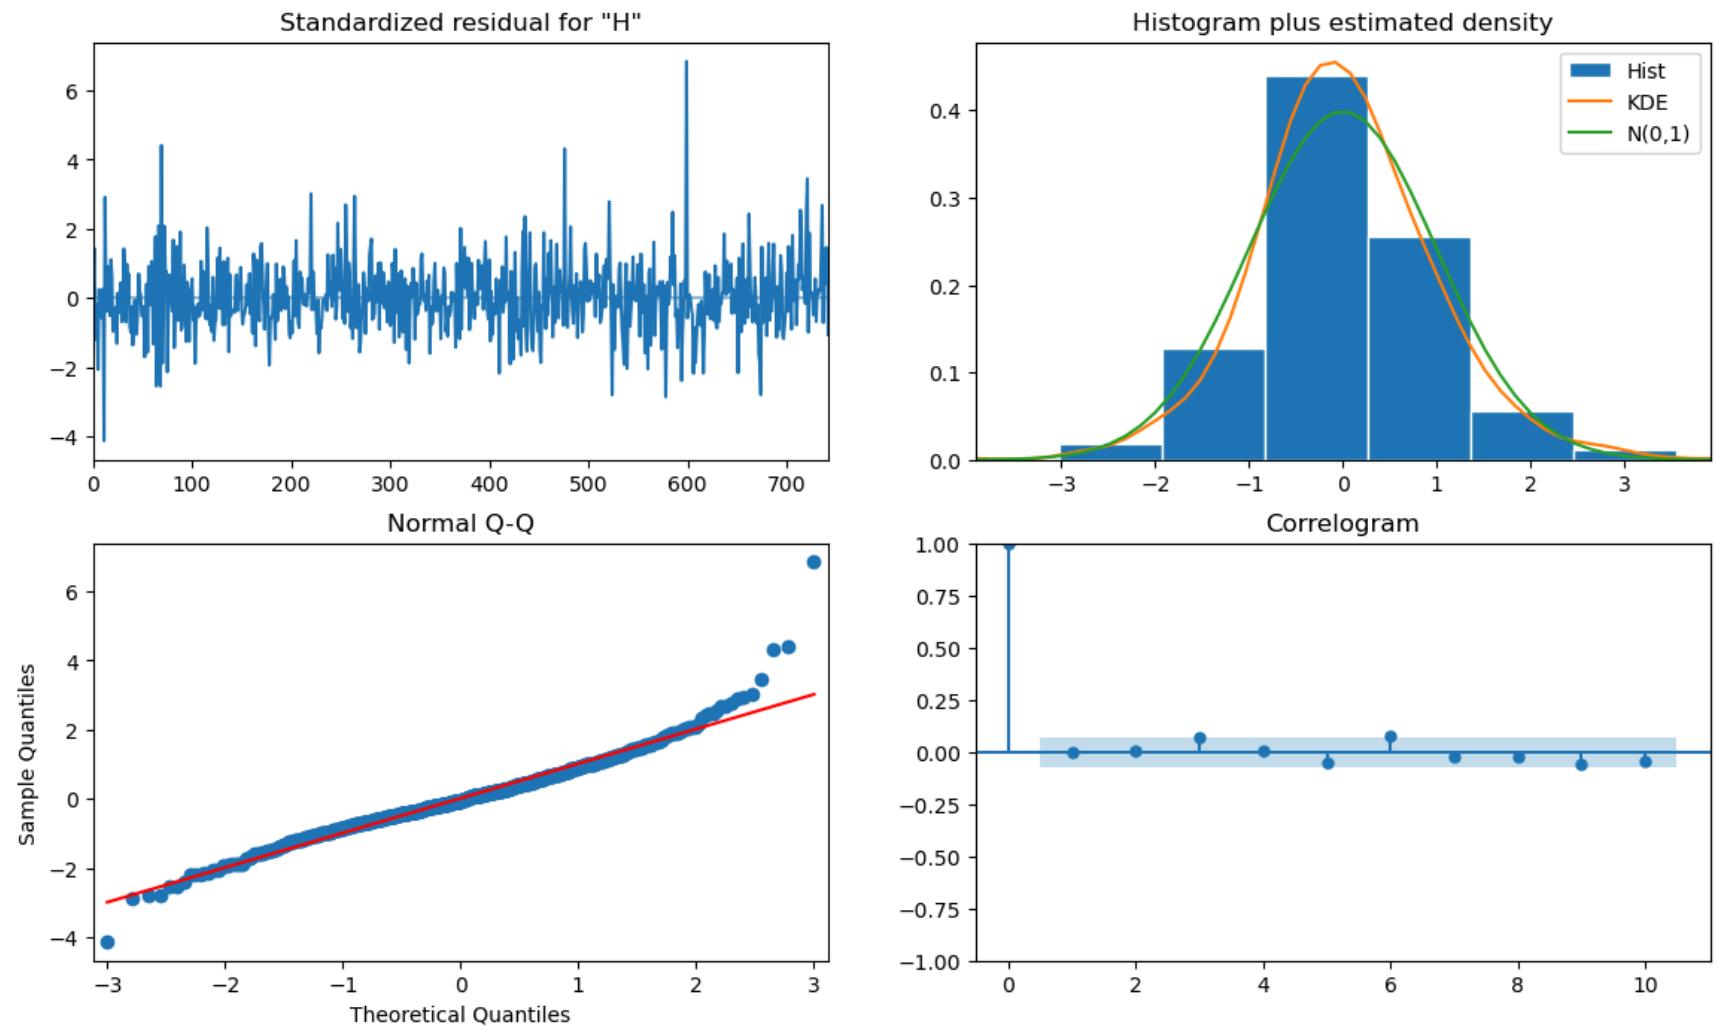
Evaluating GARNELEN

Best model for GARNELEN: $p=1, d=0, q=1, P=1, D=0, Q=1$



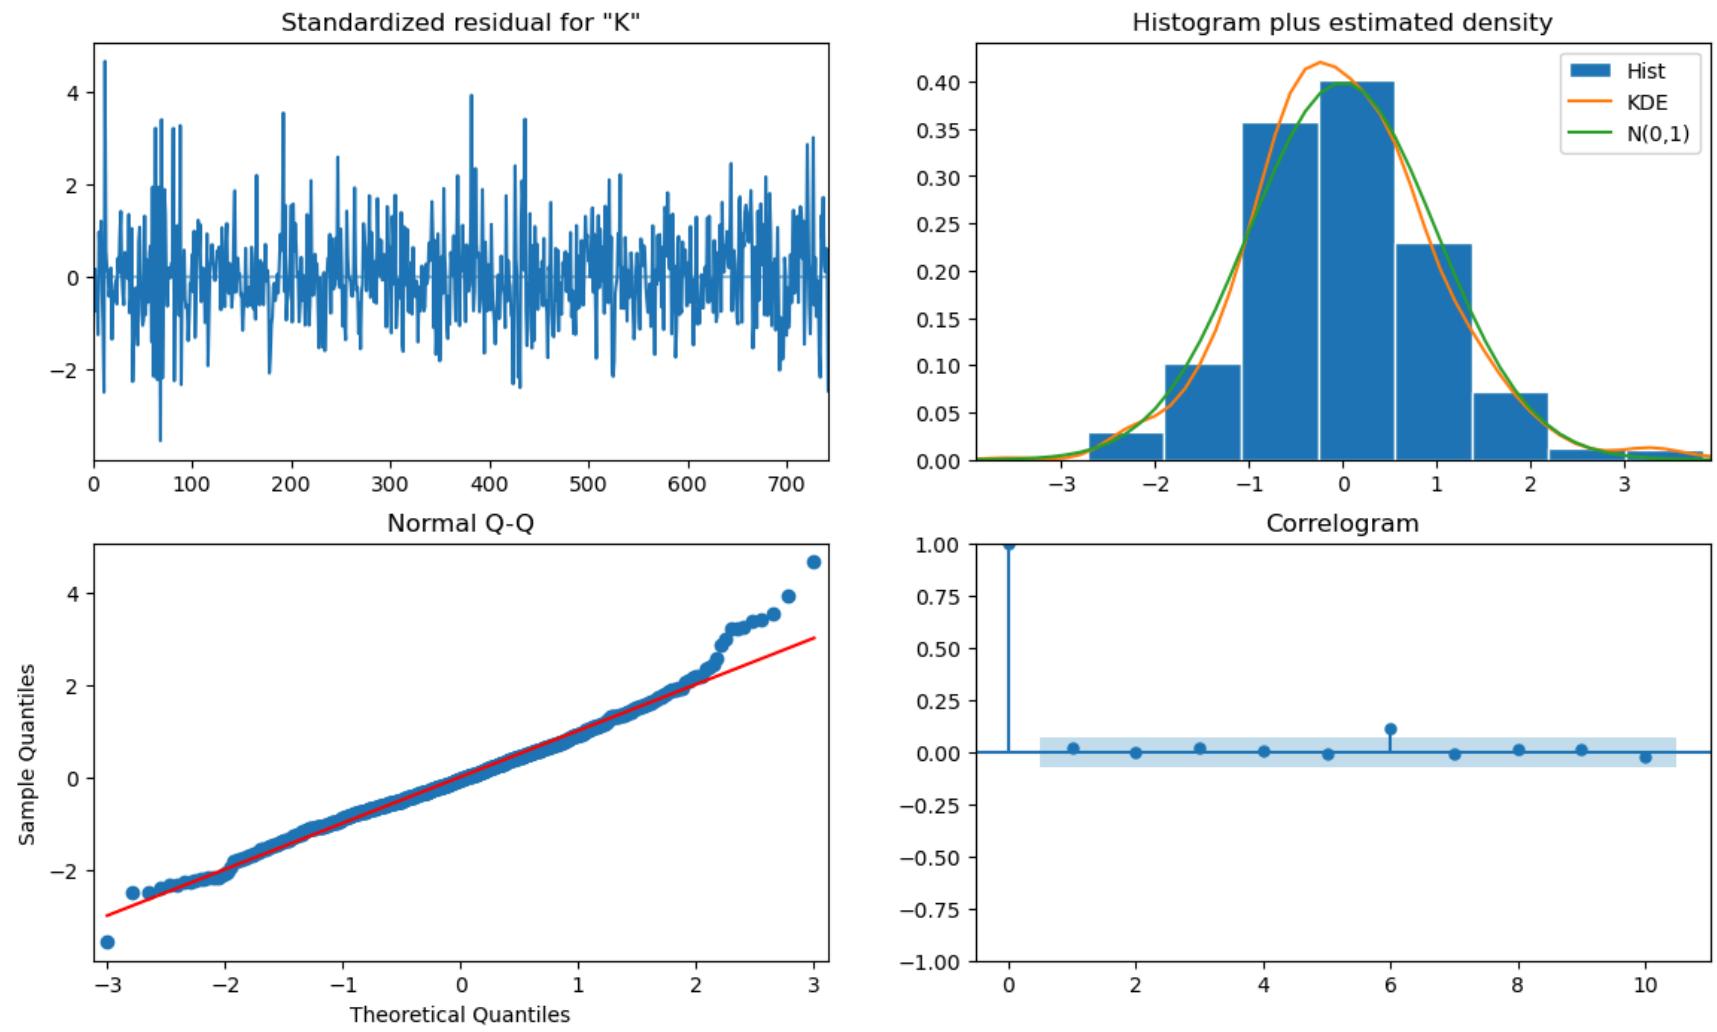
Evaluating HAEHNCHEN

Best model for HAEHNCHEN: $p=0, d=1, q=1, P=0, D=1, Q=1$



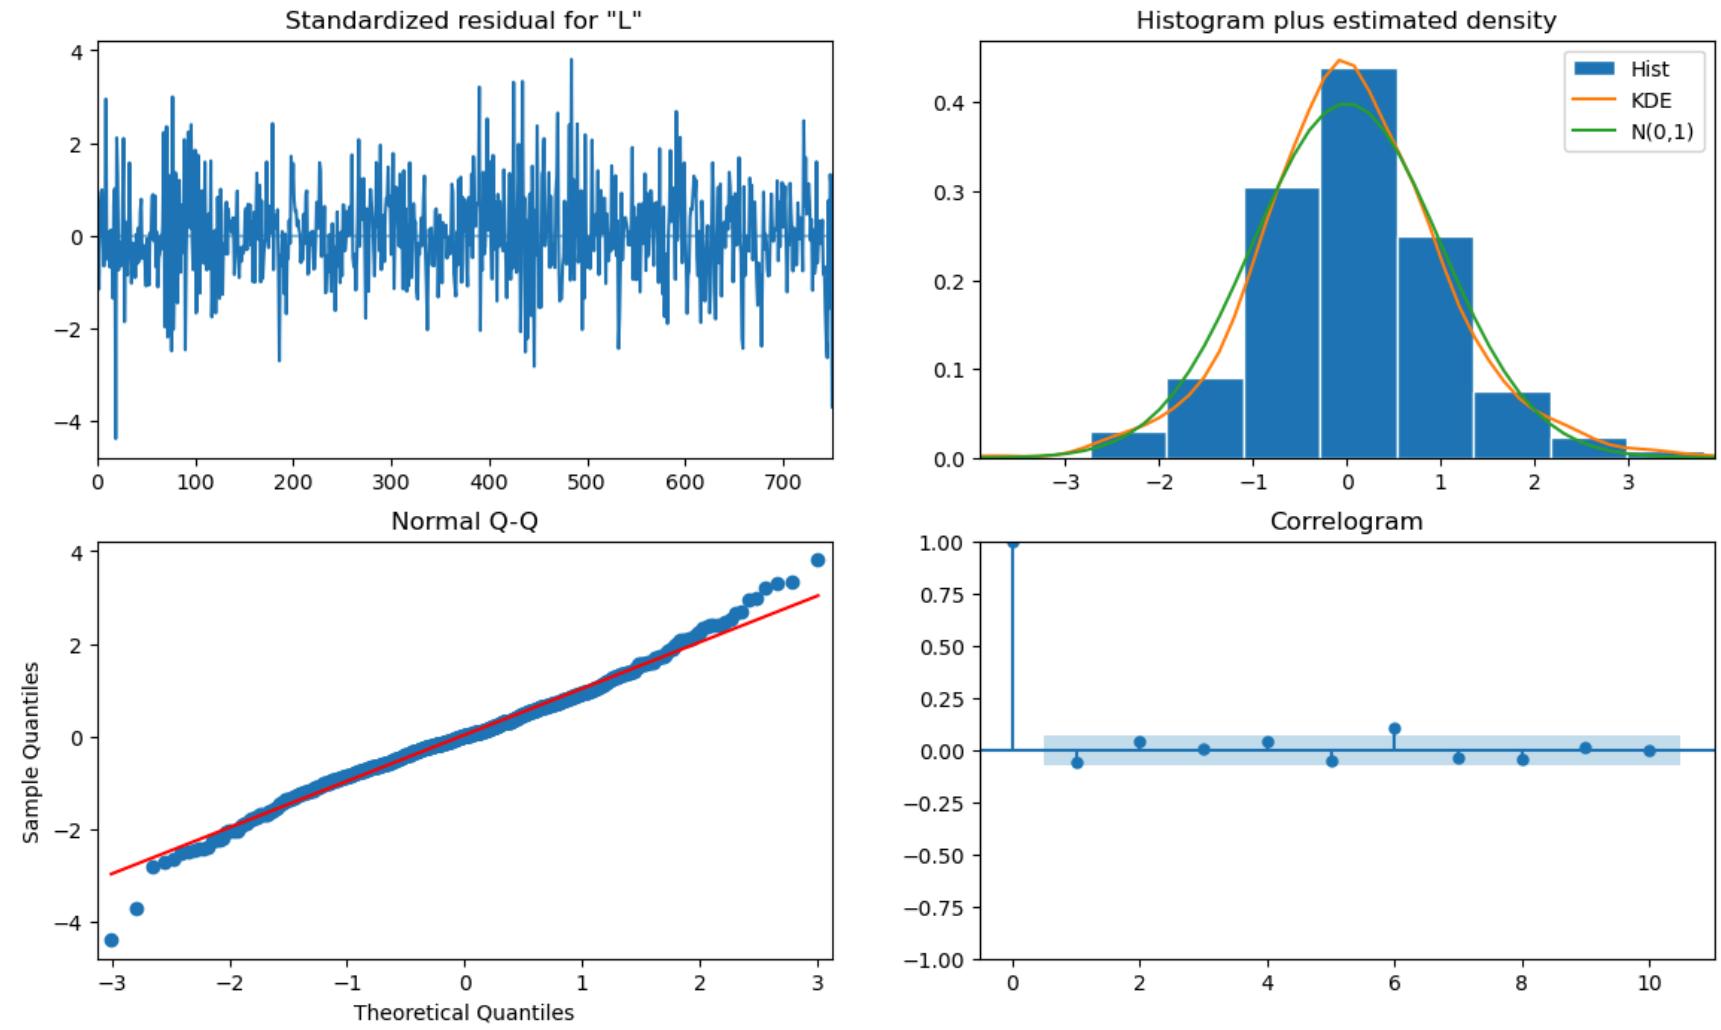
Evaluating KOEFTE

Best model for KOEFTE: $p=0, d=1, q=1, P=0, D=1, Q=1$



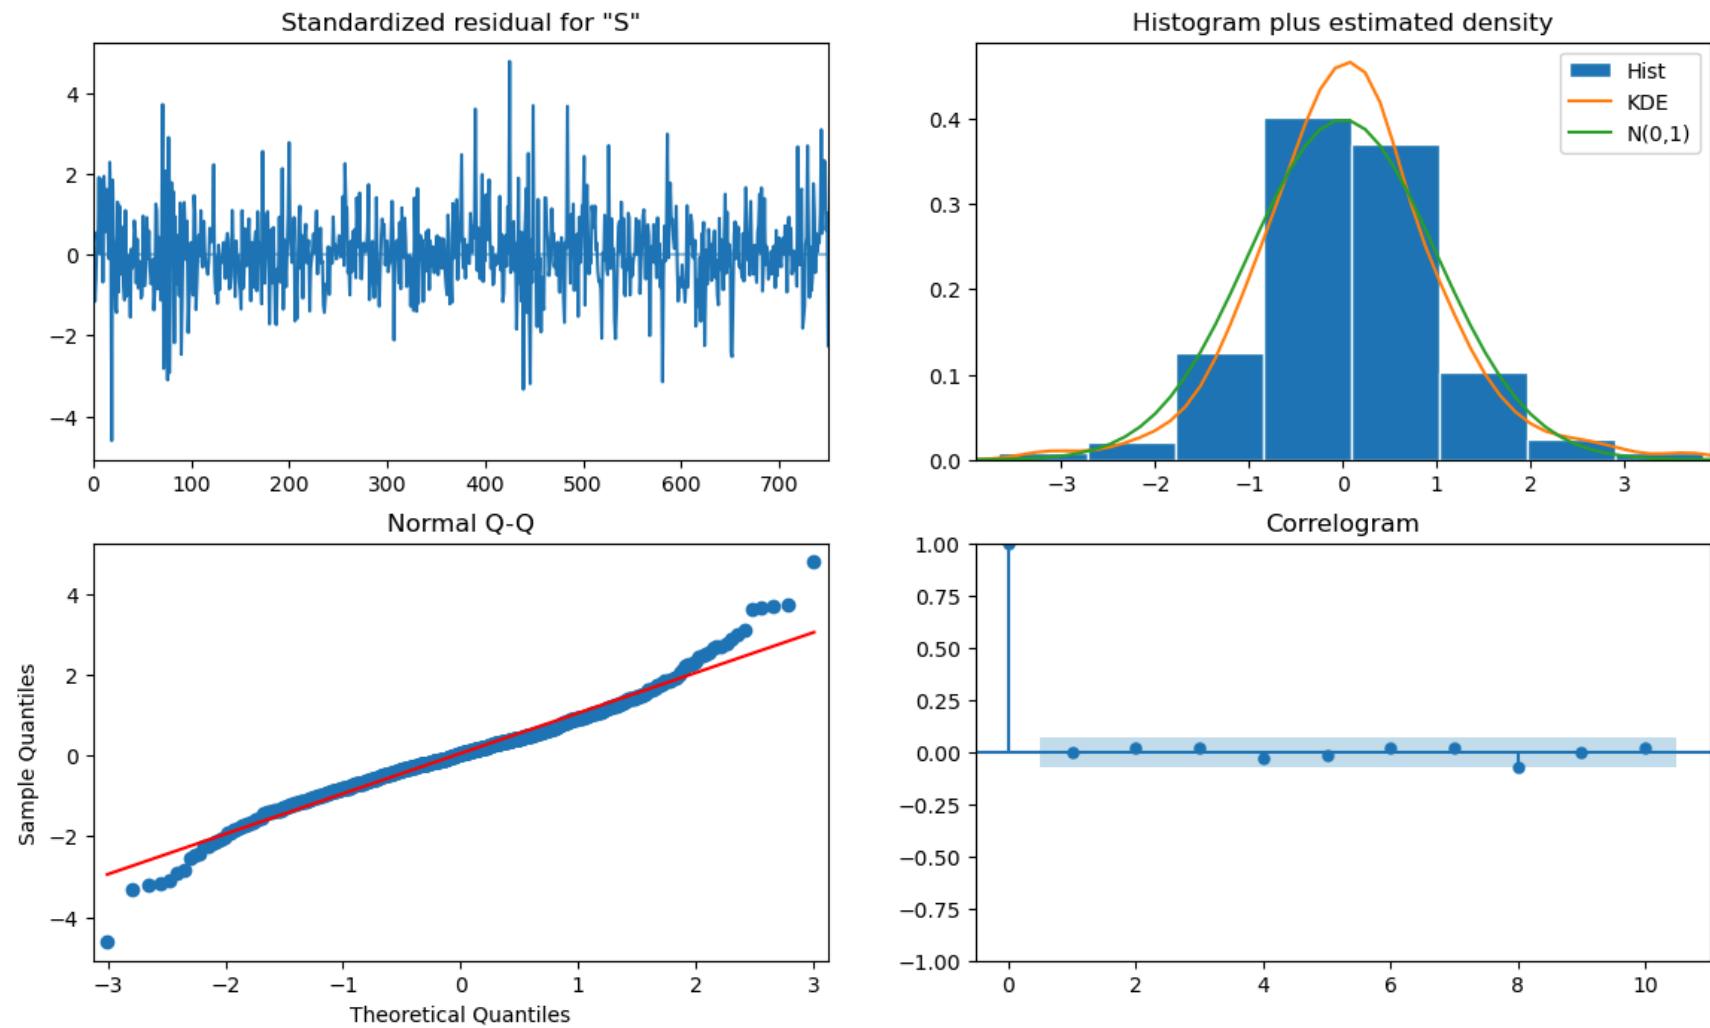
Evaluating LAMM

Best model for LAMM: $p=1, d=0, q=1, P=1, D=0, Q=1$



Evaluating STEAK

Best model for STEAK: $p=0, d=0, q=1, P=1, D=0, Q=1$



```
{'CALAMARI': {'p': 1, 'd': 0, 'q': 1, 'P': 1, 'D': 0, 'Q': 1}, 'FISCH': {'p': 0, 'd': 0, 'q': 1, 'P': 1, 'D': 0, 'Q': 1}, 'GARNELEN': {'p': 1, 'd': 0, 'q': 1, 'P': 1, 'D': 0, 'Q': 1}, 'HAEHNCHEN': {'p': 0, 'd': 1, 'q': 1, 'P': 0, 'D': 1, 'Q': 1}, 'KOEFTETE': {'p': 0, 'd': 1, 'q': 1, 'P': 0, 'D': 1, 'Q': 1}, 'LAMM': {'p': 1, 'd': 0, 'q': 1, 'P': 1, 'D': 0, 'Q': 1}}
```

To analyze the diagnostic plots that were outputted, we check for:

Standardized Residuals: Residuals should behave like white noise (i.e., no obvious patterns, constant variance, and mean around zero). Any patterns or trends may suggest the model is missing some information.

Histogram plus KDE (Kernel Density Estimate): This plot should follow a normal bell curve.

Normal Q-Q Plot: Points should fall approximately along the 45-degree line. Significant deviations indicate that the residuals are not normally distributed, which may affect the validity of confidence intervals and forecasts.

Correlogram of the Residuals (ACF Plot): This plot shows the autocorrelation of the residuals at various lags. There should not be any significant spikes outside the shaded confidence interval area.

The plots that were outputted for each menu item after finding the best parameters fulfill all the requirements that we are looking for in the plot diagnostics to make sure that the model has a quality fit.

Model Evaluation

```
In [ ]: def evaluate_sarima(item, train, test, test_size, params, m):

    model = SARIMAX(
        df_time_series[item],
        order=(params[item]['p'], params[item]['d'], params[item]['q']),
        seasonal_order=(params[item]['P'], params[item]['D'], params[item]['Q'], m),
        enforce_stationarity=False,
        enforce_invertibility=False
    )

    sarima_model = model.fit(disp=False)

    forecast = sarima_model.get_forecast(steps=test_size)
    forecast_index = test.index
    forecast_values = forecast.predicted_mean

    plt.figure(figsize=(12, 6))
    plt.plot(test.index, test[item], label='Test', color='green')
    plt.plot(forecast_index, forecast_values, label='Forecast', color='red')
    plt.title(f'{item}: SARIMA Forecast vs. Actual')
    plt.xlabel('Date')
    plt.ylabel('Sales')
    plt.legend()
    plt.grid(True)
    plt.show()

    mse = round(mean_squared_error(test[item], forecast_values), 2)
    print(f'{item} Mean Squared Error: {mse}')

return mse
```

In [327...]

```
def run(m):
    mse_all_items = {}

    for item in menu_items:
        test_size = 7 * 16 # Last 16 weeks
        train_end = len(df_time_series) - test_size
        train = df_time_series.iloc[:train_end]
        test = df_time_series.iloc[train_end:]

        mse = evaluate_sarima(item, train, test, test_size, best_params, m)
        mse_all_items[item] = mse

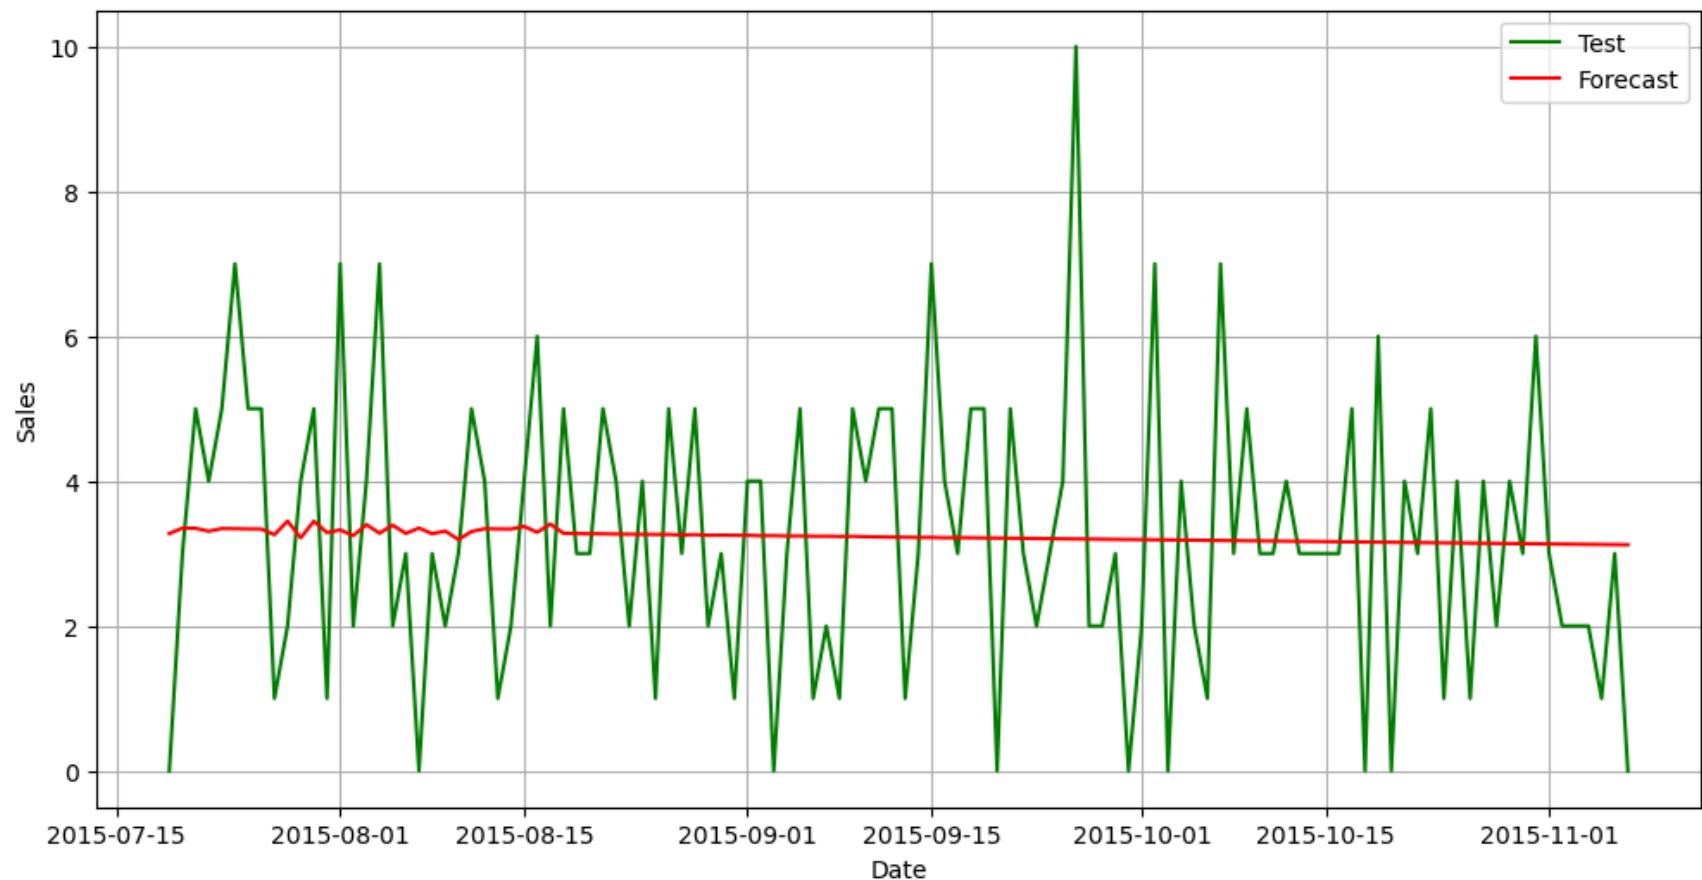
    return mse_all_items
```

Comparing Monthly Trends (m = 30) vs. Weekly Trends (m = 7)

In [328...]

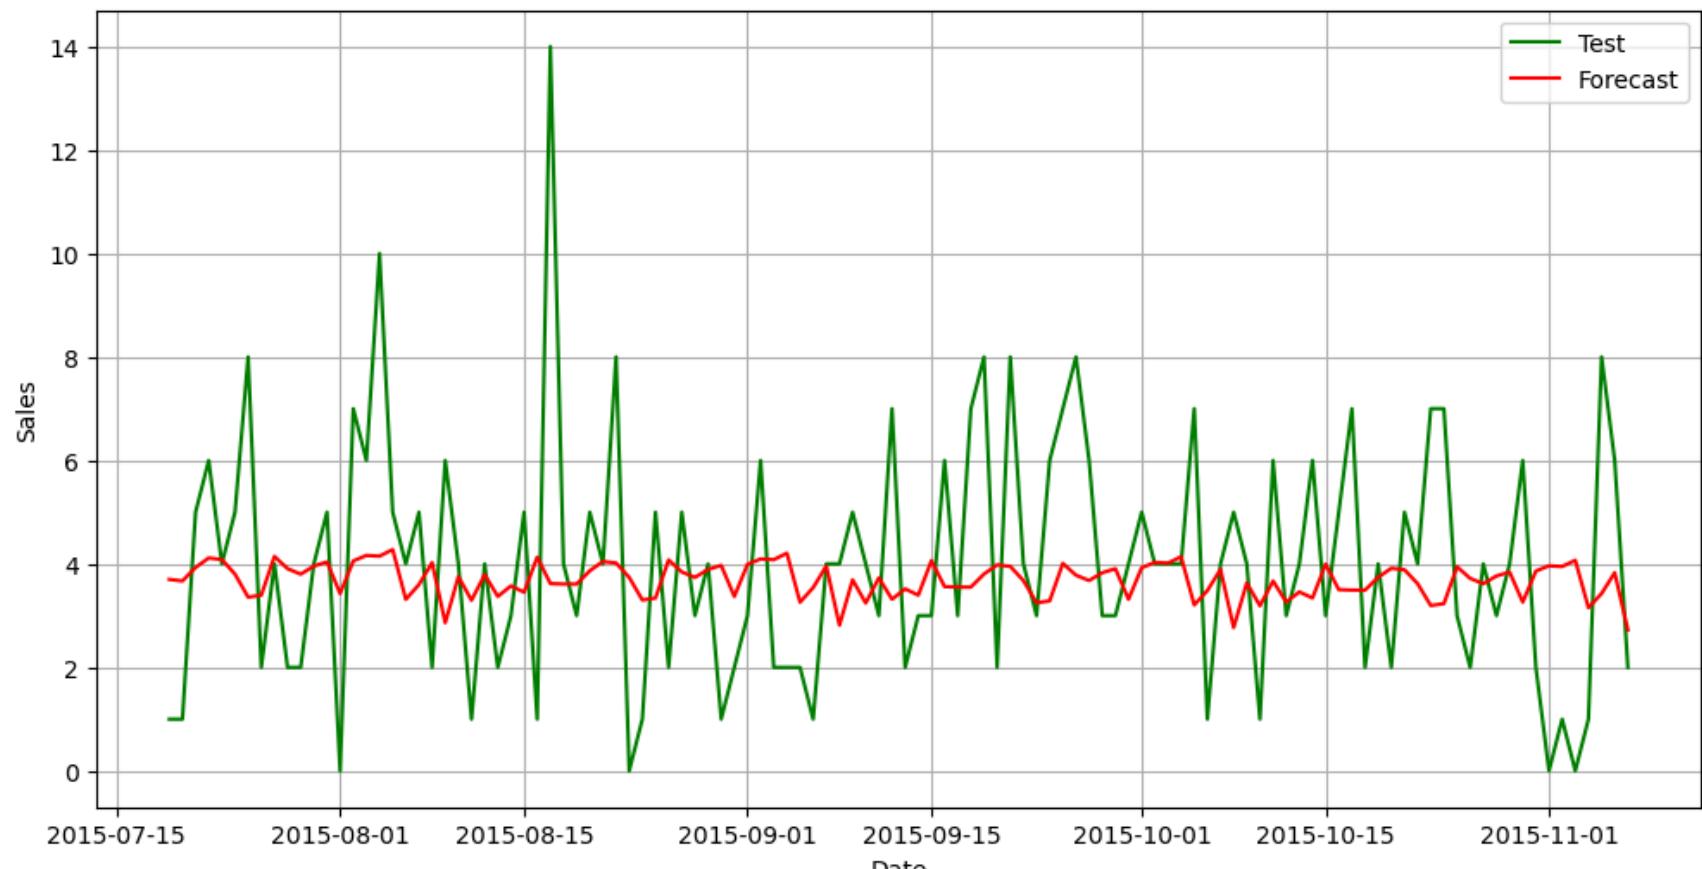
```
run(m=30)
```

CALAMARI: SARIMA Forecast vs. Actual



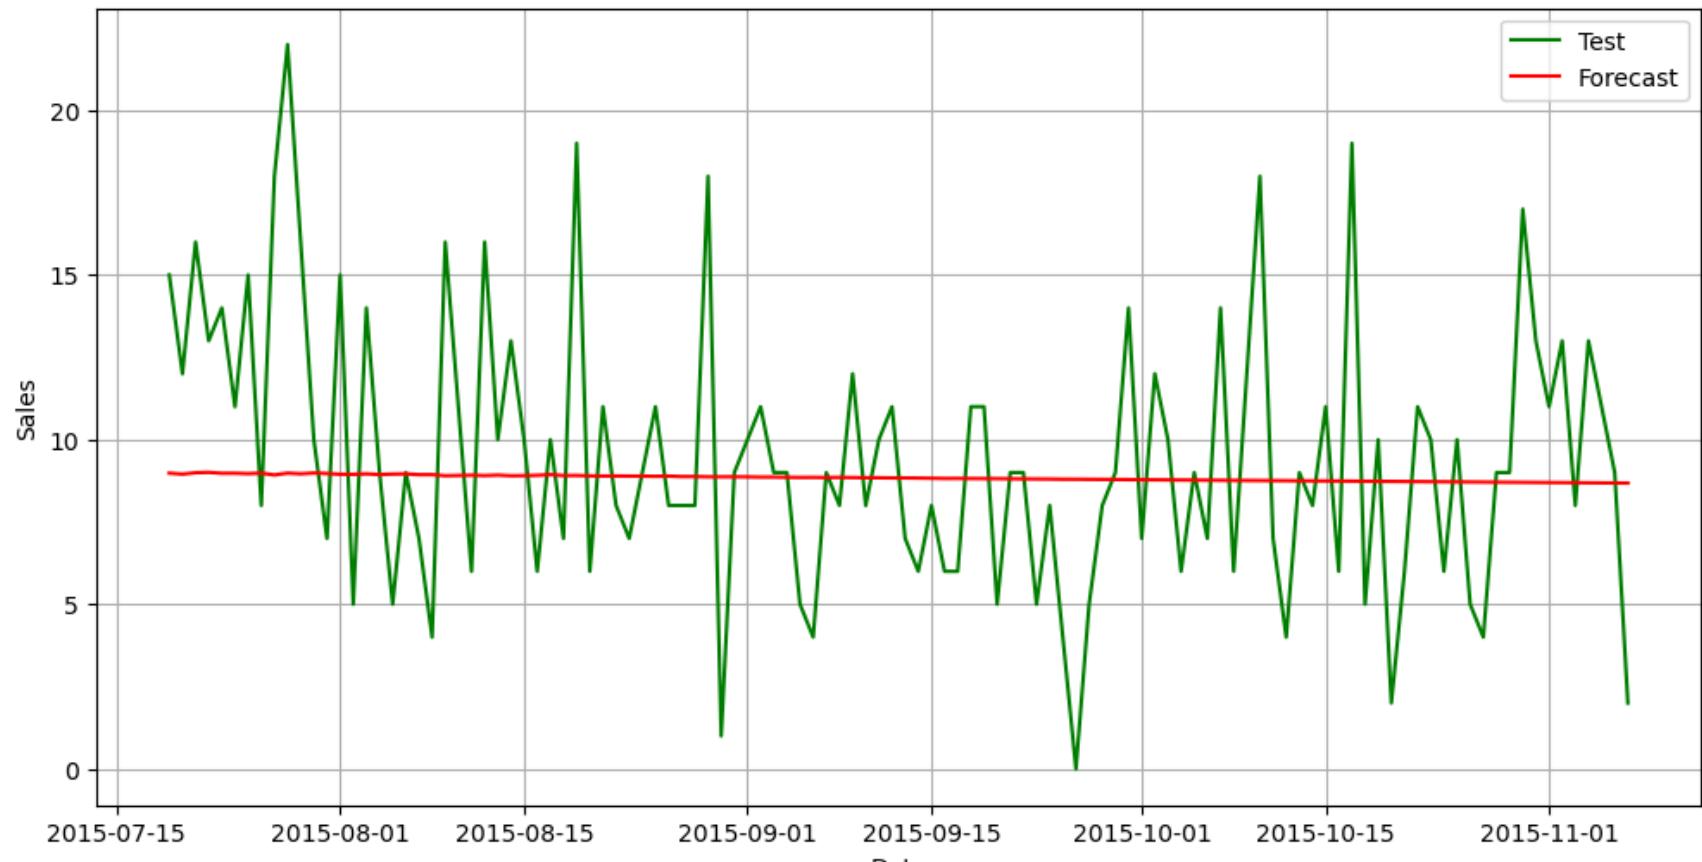
CALAMARI Mean Squared Error: 3.63

FISCH: SARIMA Forecast vs. Actual



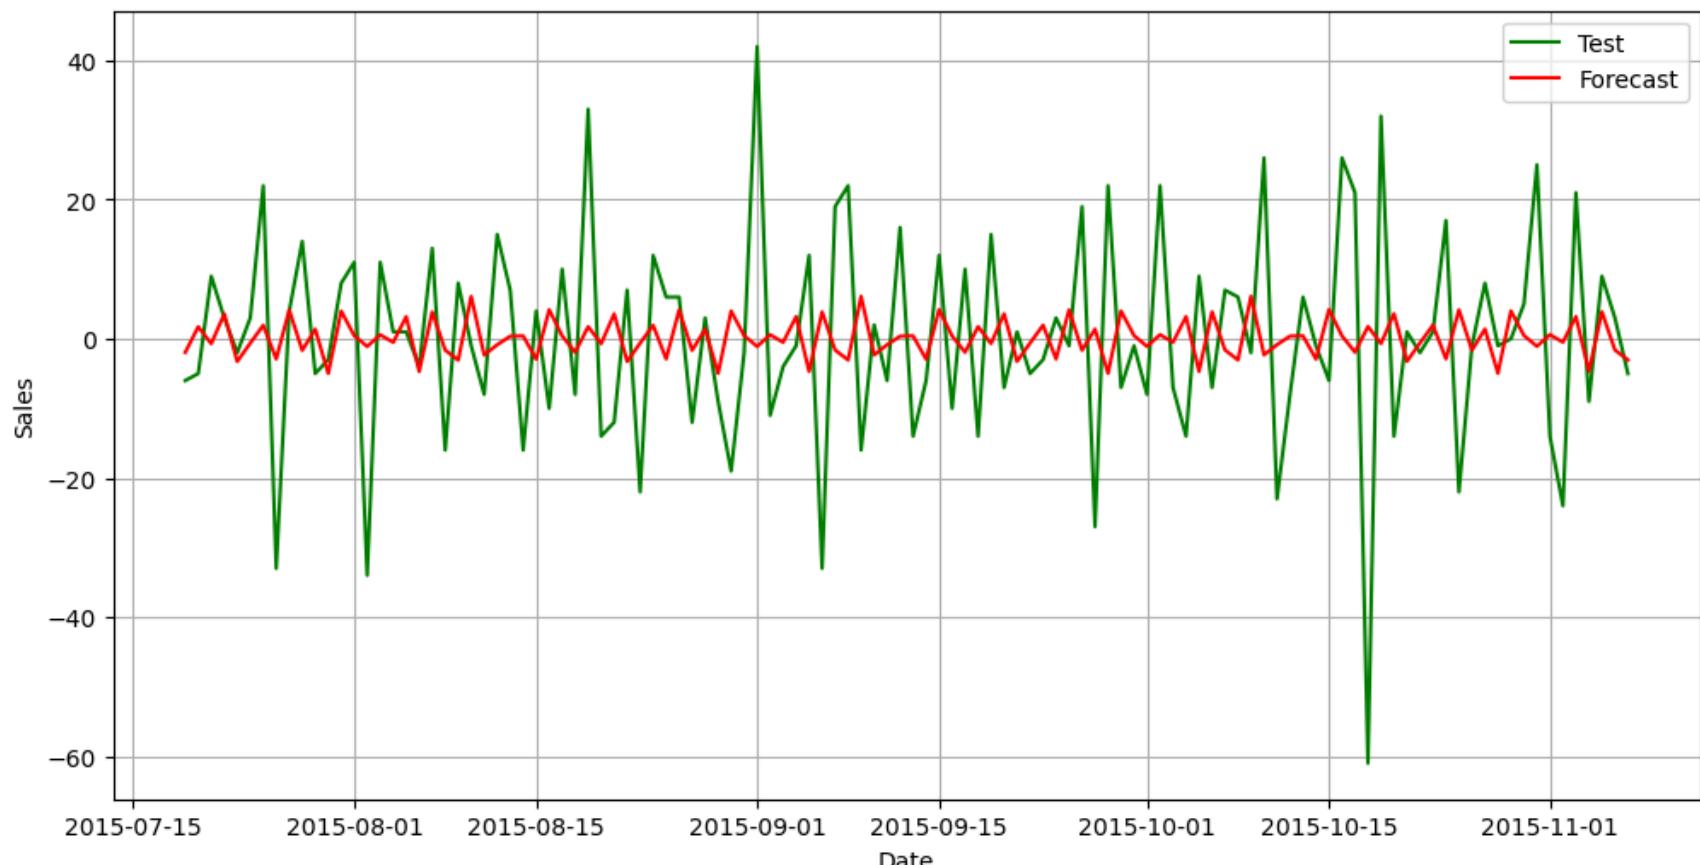
FISCH Mean Squared Error: 5.67

GARNELEN: SARIMA Forecast vs. Actual



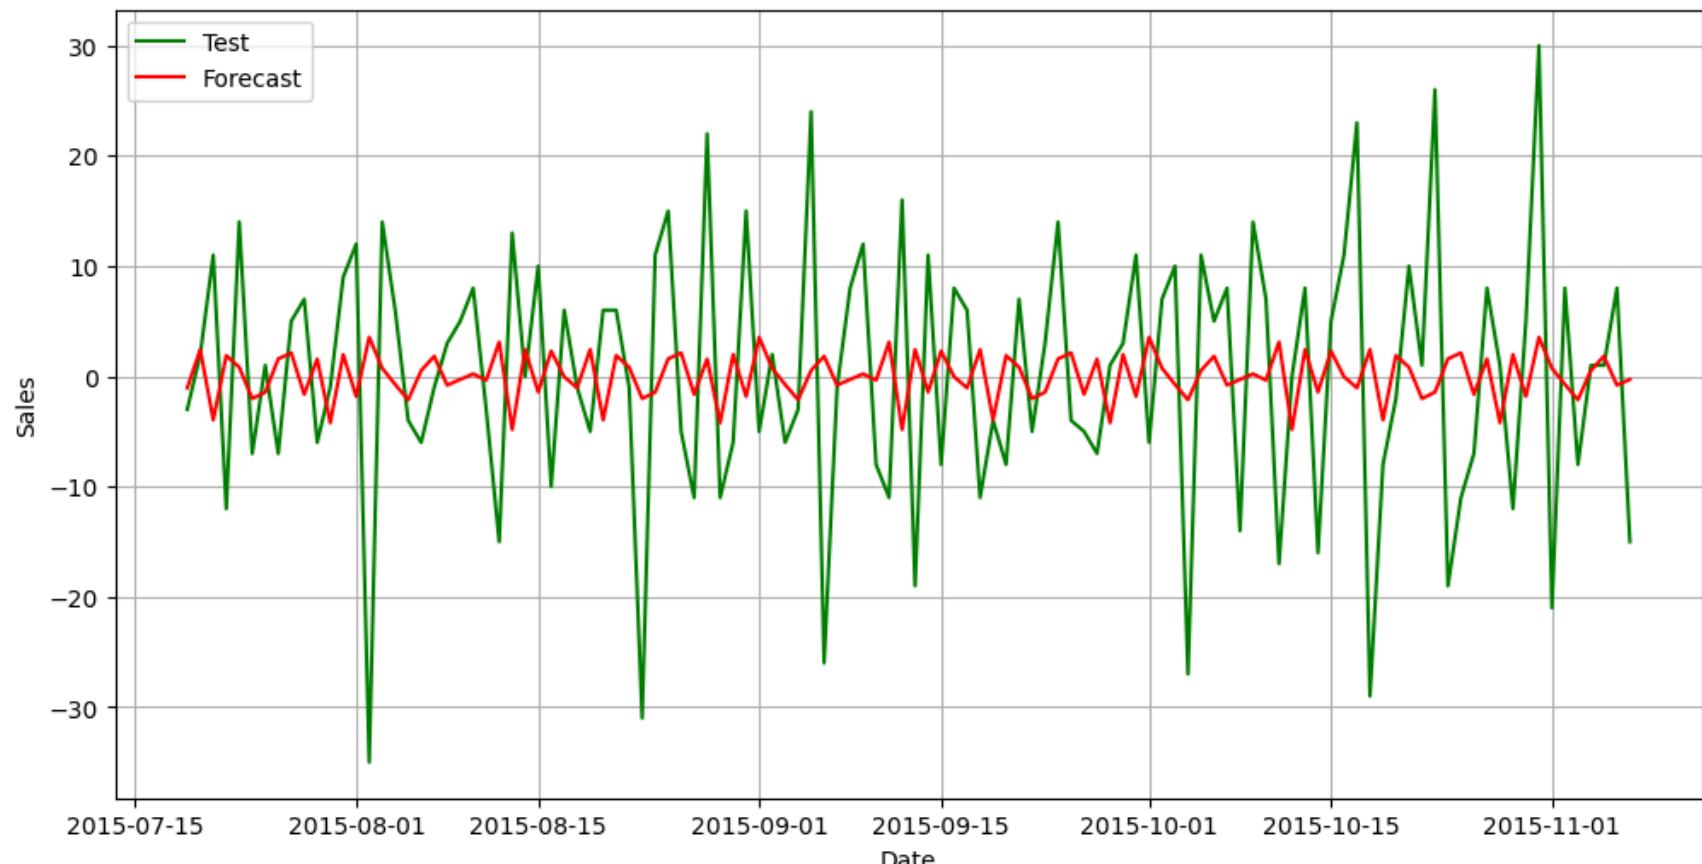
GARNELEN Mean Squared Error: 16.92

HAEHNCHEN: SARIMA Forecast vs. Actual



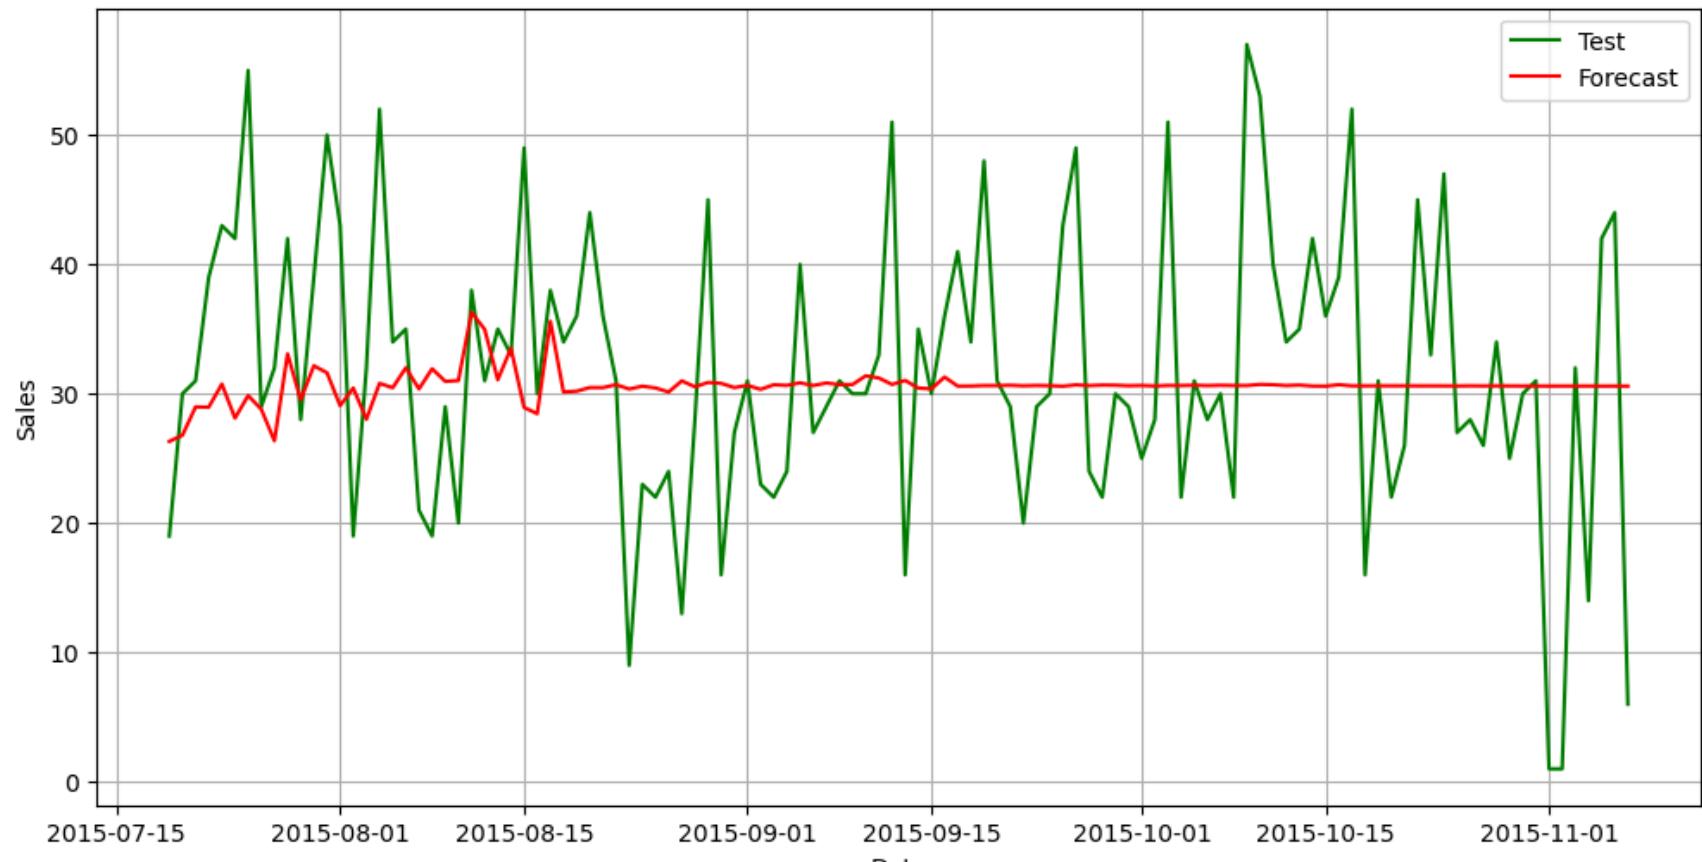
HAEHNCHEN Mean Squared Error: 255.04

KOEFTE: SARIMA Forecast vs. Actual



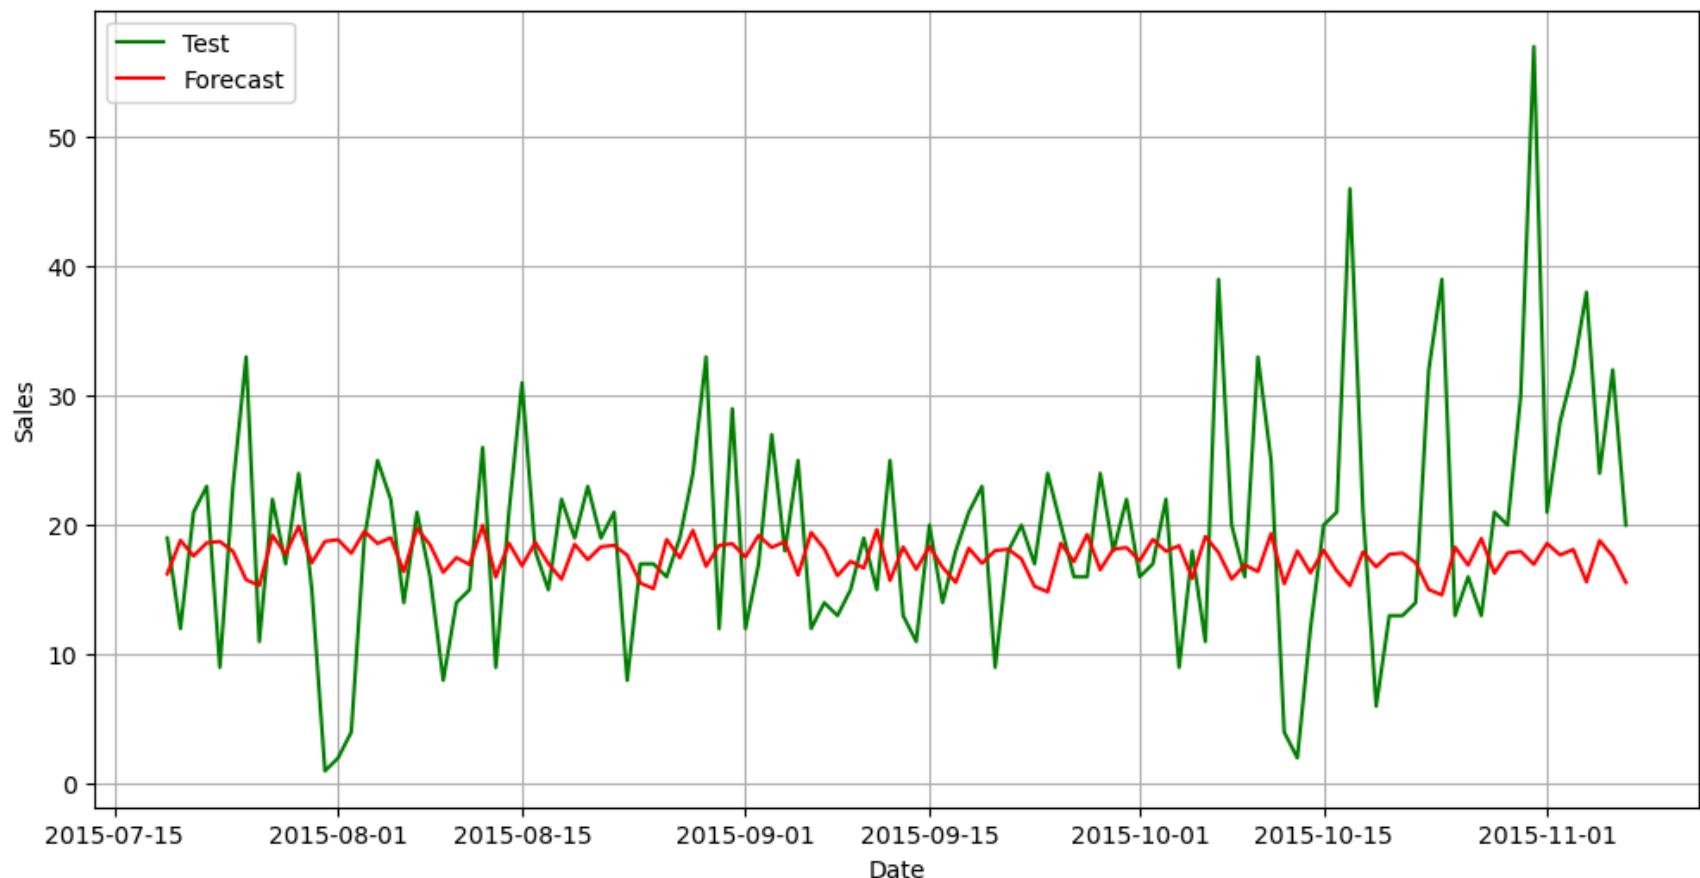
KOEFTTE Mean Squared Error: 155.85

LAMM: SARIMA Forecast vs. Actual



LAMM Mean Squared Error: 120.16

STEAK: SARIMA Forecast vs. Actual



STEAK Mean Squared Error: 84.47

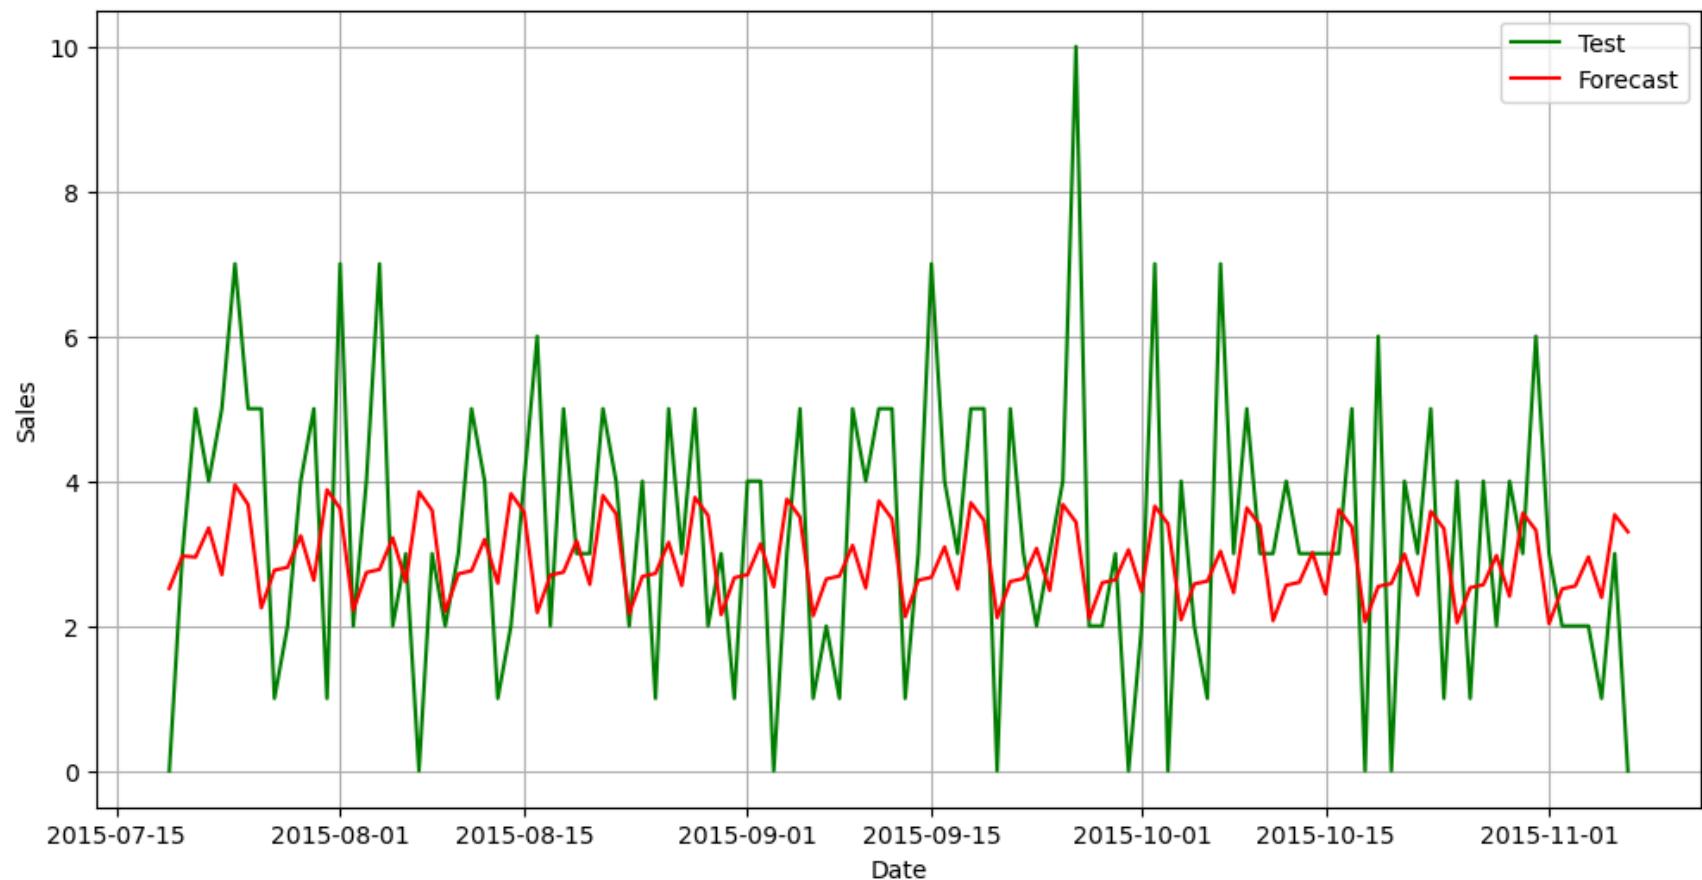
Out[328]:

```
{'CALAMARI': 3.63,
 'FISCH': 5.67,
 'GARNELEN': 16.92,
 'HAEHNCHEN': 255.04,
 'KOEFTETE': 155.85,
 'LAMM': 120.16,
 'STEAK': 84.47}
```

The monthly seasonal period does not seem to perform very well, as it does not correctly capture the trends that the data exhibits. We can see this in LAMM for example, as the model becomes stagnant at the end of the month, ignoring the dips and peaks in the data.

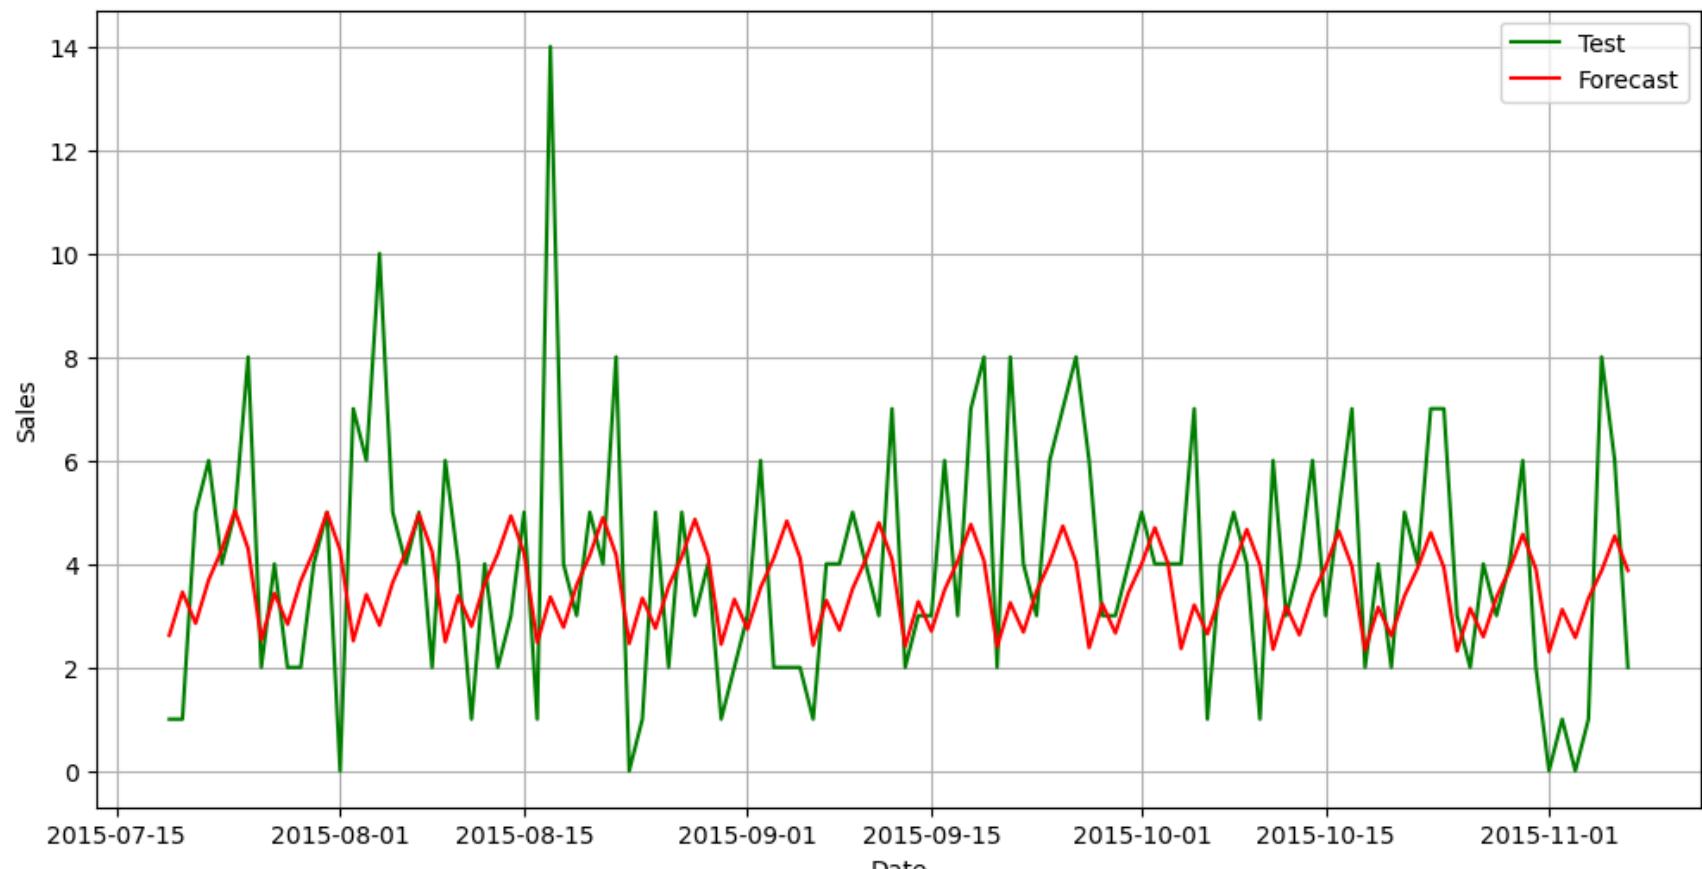
In [329... run(m=7)

CALAMARI: SARIMA Forecast vs. Actual



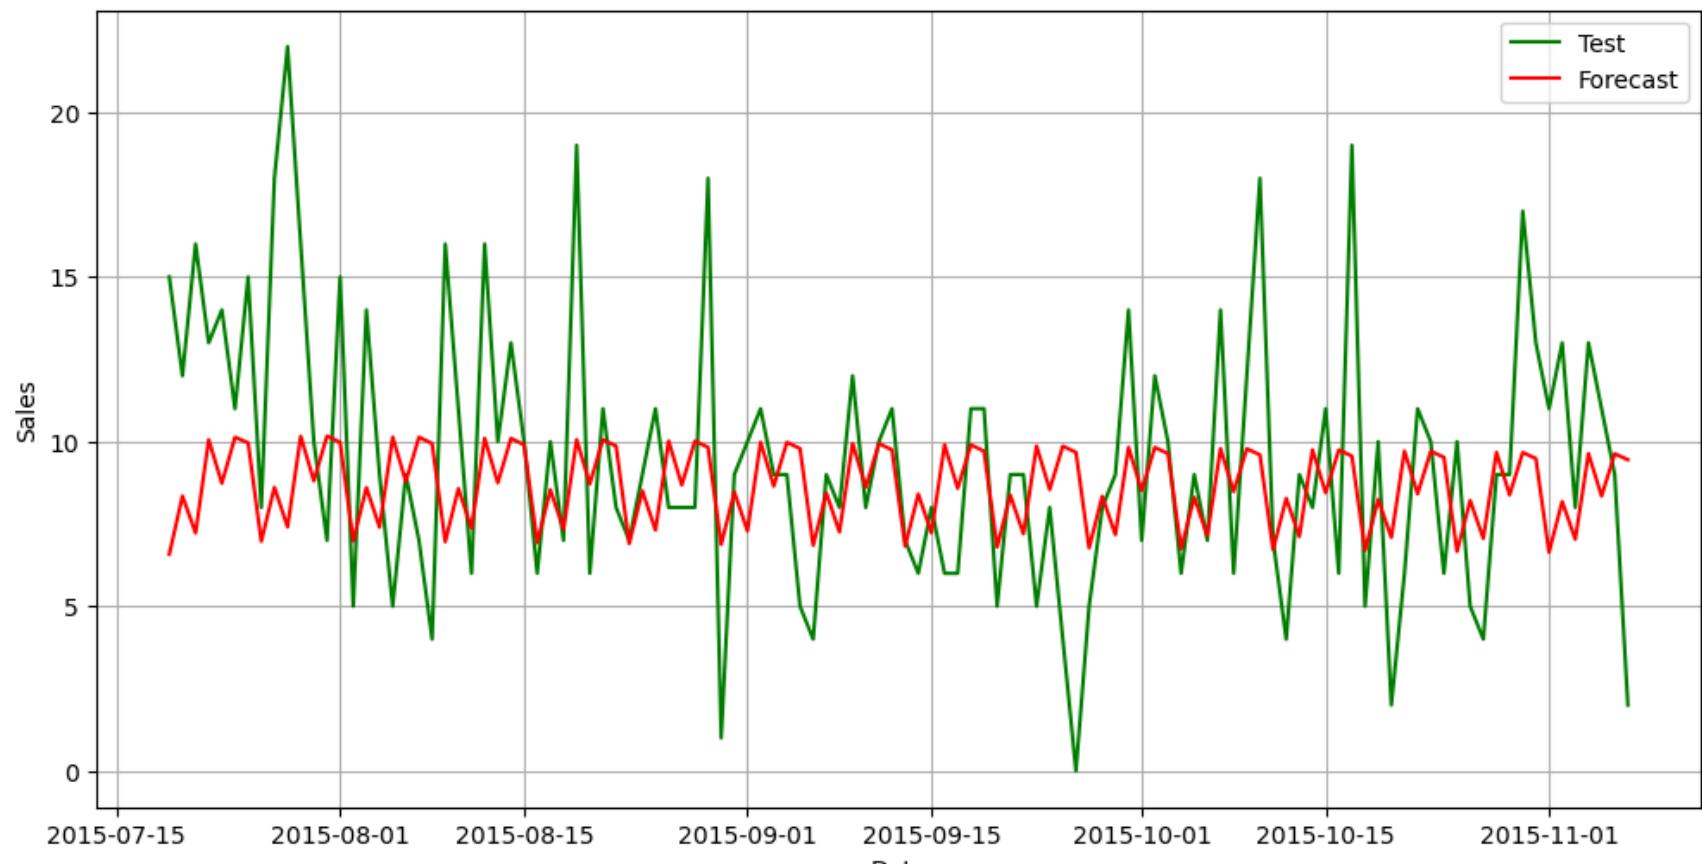
CALAMARI Mean Squared Error: 3.45

FISCH: SARIMA Forecast vs. Actual



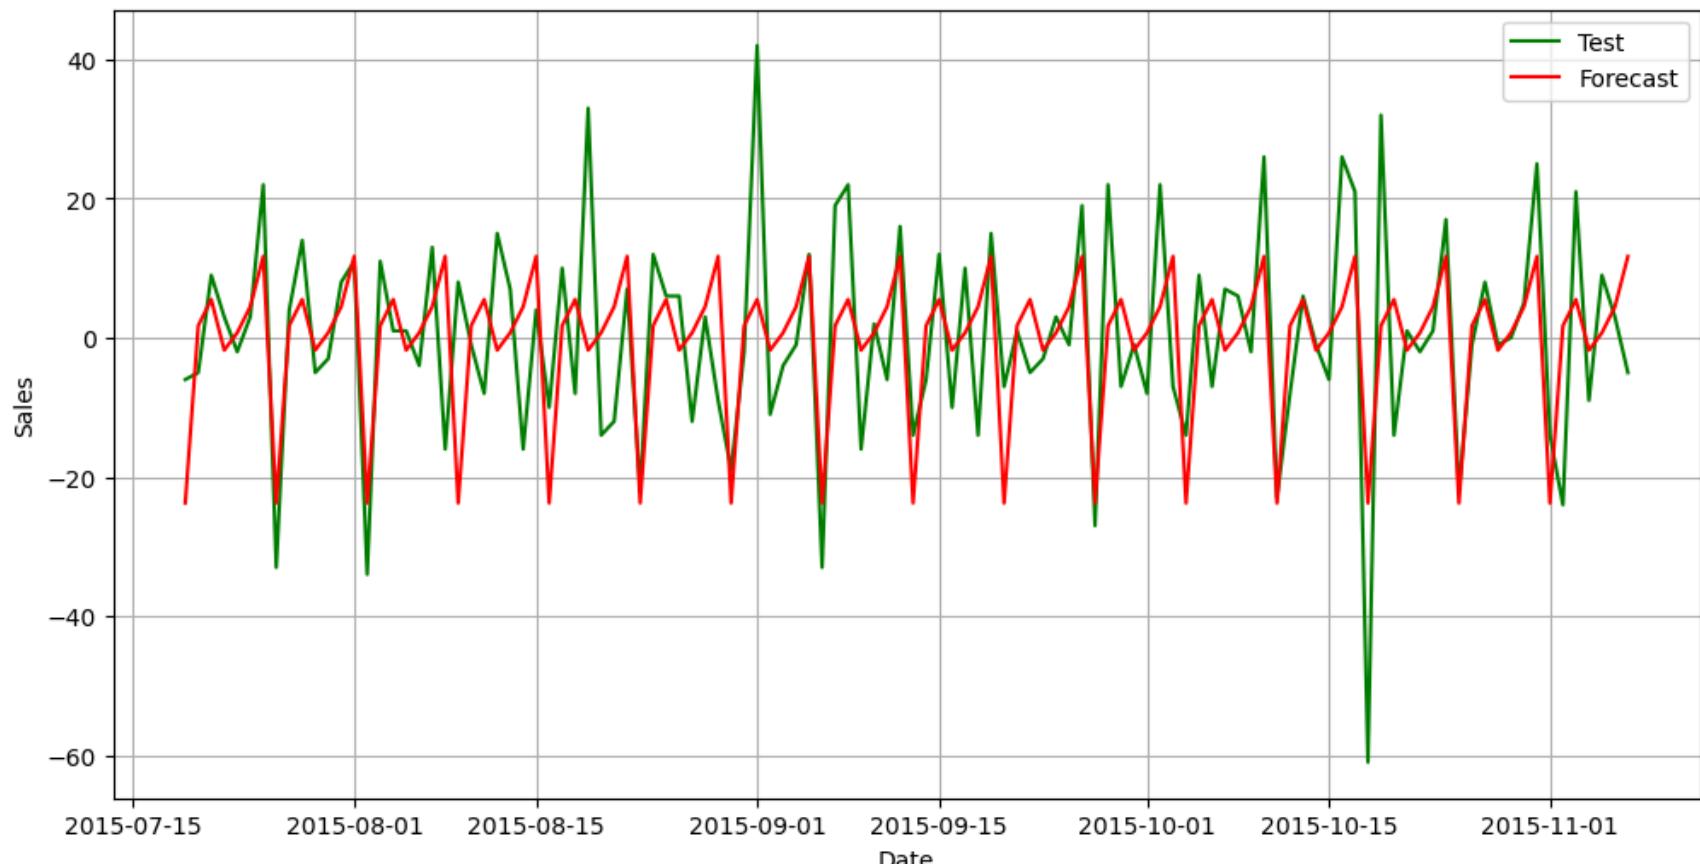
FISCH Mean Squared Error: 5.32

GARNELEN: SARIMA Forecast vs. Actual



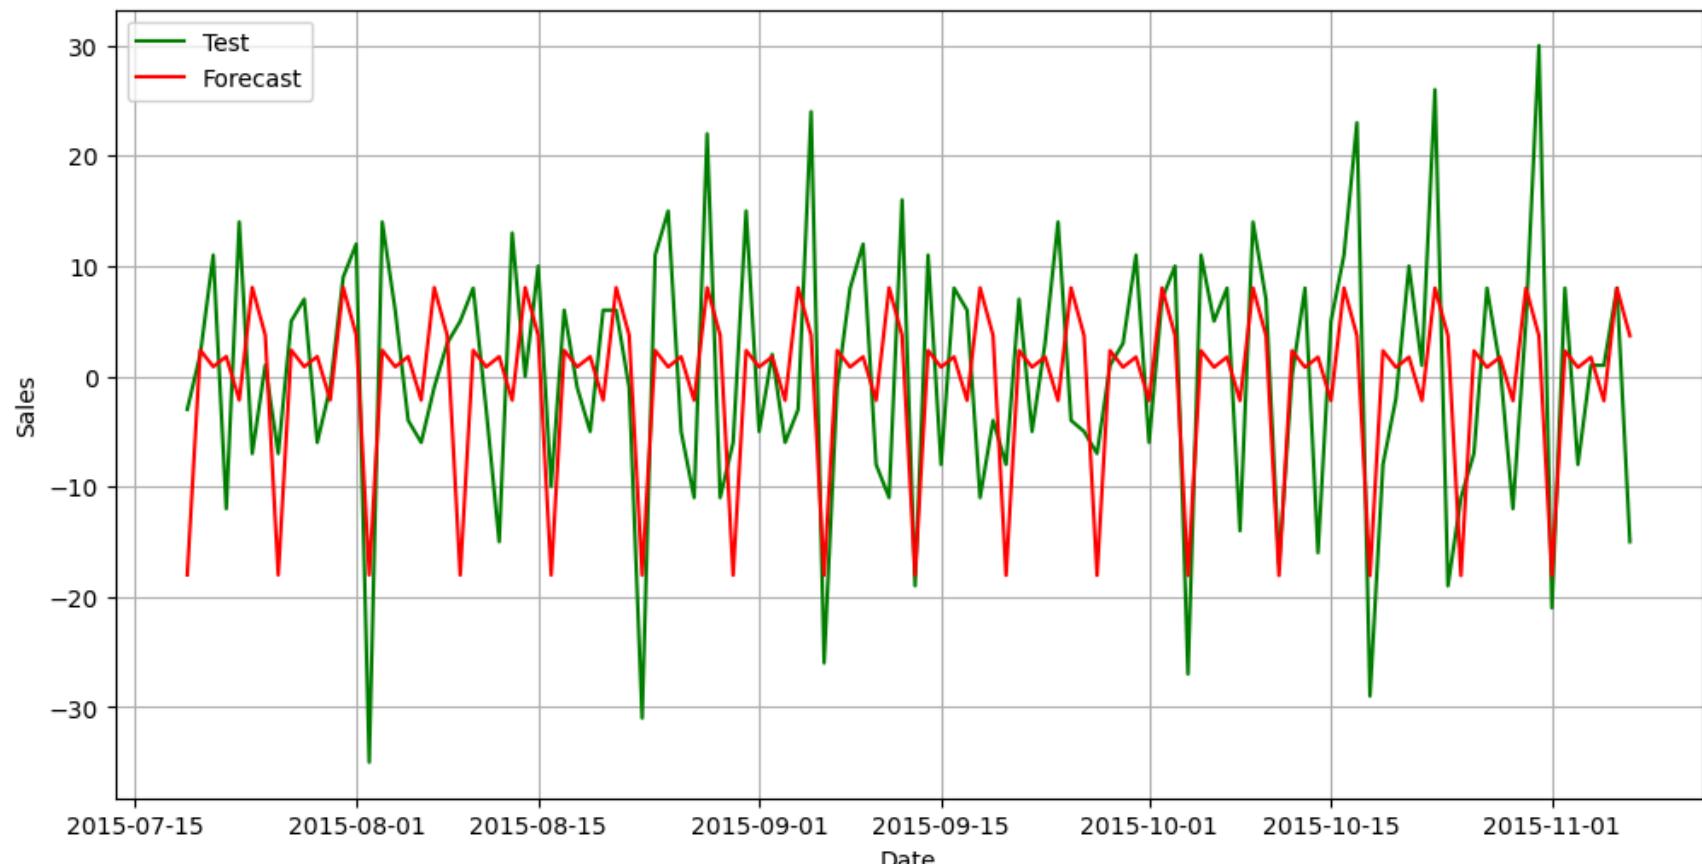
GARNELEN Mean Squared Error: 16.44

HAEHNCHEN: SARIMA Forecast vs. Actual



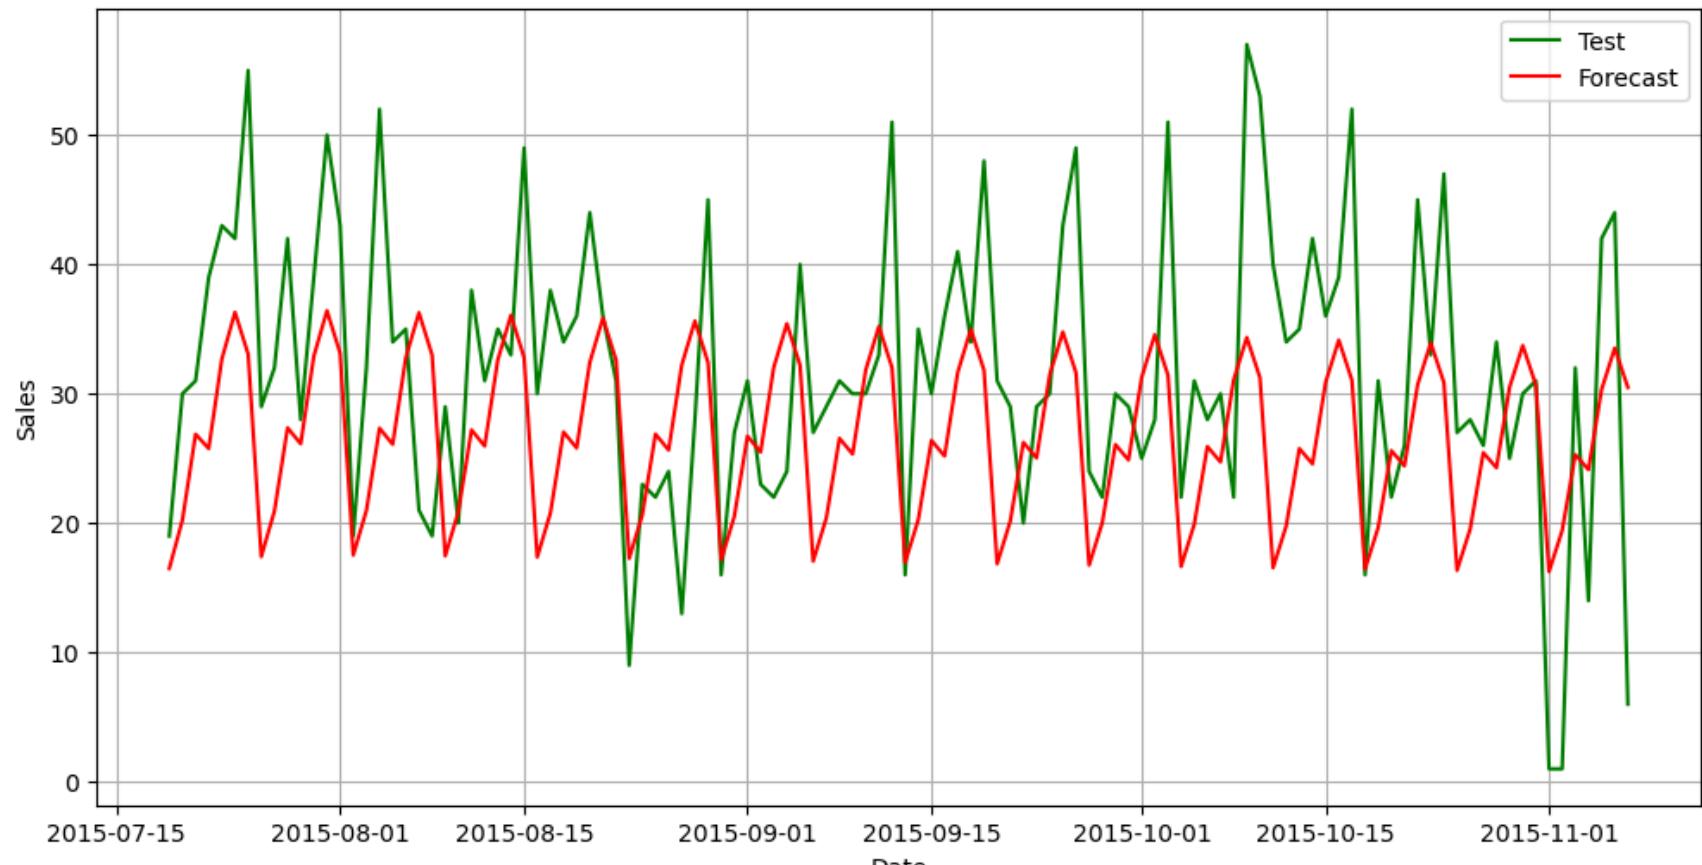
HAEHNCHEN Mean Squared Error: 155.18

KOEFTE: SARIMA Forecast vs. Actual



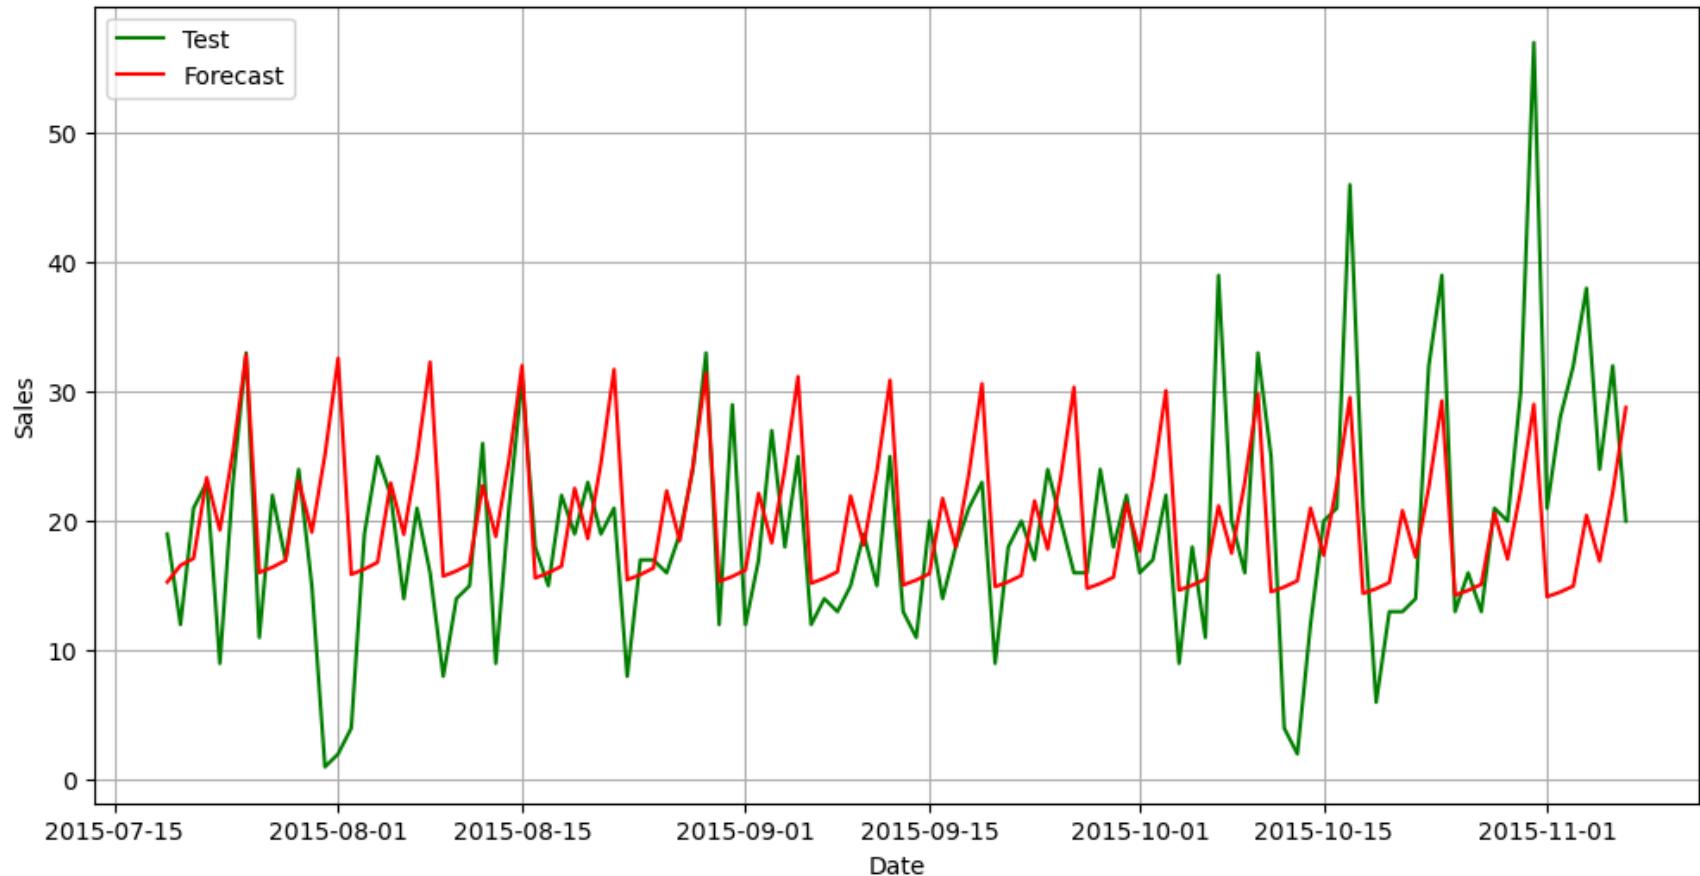
KOEFTE Mean Squared Error: 101.42

LAMM: SARIMA Forecast vs. Actual



LAMM Mean Squared Error: 117.89

STEAK: SARIMA Forecast vs. Actual



STEAK Mean Squared Error: 67.3

Out[329]:

```
{'CALAMARI': 3.45,
 'FISCH': 5.32,
 'GARNELEN': 16.44,
 'HAEHNCHEN': 155.18,
 'KOEFTETE': 101.42,
 'LAMM': 117.89,
 'STEAK': 67.3}
```

The visualization for a weekly seasonal period fits the trends in the data much better than a monthly seasonal period. We can see that generally the spikes follow the natural dips and peaks of weekly sales (higher sales on weekends, followed by a decrease in sales on the weekdays).

In [506...]

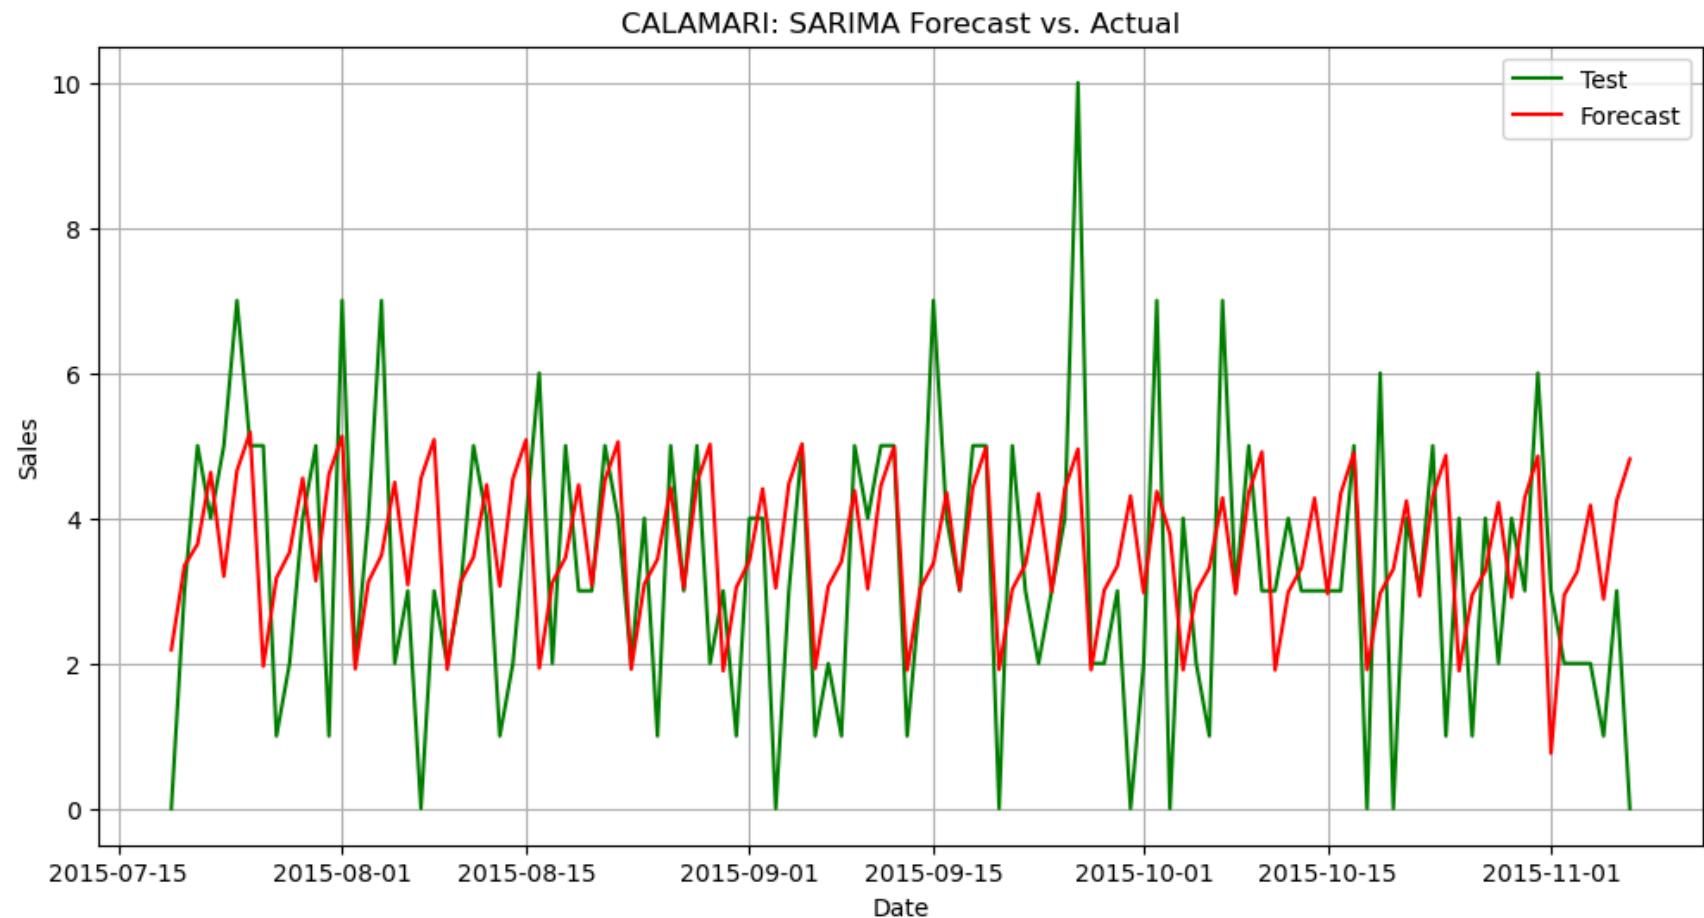
```
def evaluate_sarima_exog(item, train, test, test_size, params, m, train_exog, test_exog):  
    model = SARIMAX(  
        train[item],  
        order=(params[item]['p'], params[item]['d'], params[item]['q']),  
        seasonal_order=(params[item]['P'], params[item]['D'], params[item]['Q'], m),  
        exog=train_exog,  
        enforce_stationarity=False,  
        enforce_invertibility=False  
    )  
  
    sarima_model = model.fit(disp=False)  
  
    forecast = sarima_model.get_forecast(steps=test_size, exog=test_exog)  
    forecast_index = test.index  
    forecast_values = forecast.predicted_mean  
  
    plt.figure(figsize=(12, 6))  
    plt.plot(test.index, test[item], label='Test', color='green')  
    plt.plot(forecast_index, forecast_values, label='Forecast', color='red')  
    plt.title(f'{item}: SARIMA Forecast vs. Actual')  
    plt.xlabel('Date')  
    plt.ylabel('Sales')  
    plt.legend()  
    plt.grid(True)  
    plt.show()  
  
    mse = round(mean_squared_error(test[item], forecast_values), 2)  
    print(f'{item} Mean Squared Error: {mse}')  
  
    return mse  
  
def run_exog(m):  
    mse_all_items = []  
  
    for item in menu_items:  
        test_size = 7 * 16 # Last 16 weeks  
        train_end = len(df_time_series) - test_size  
        train = df_time_series.iloc[:train_end]  
        test = df_time_series.iloc[train_end:]  
  
        train_exog = df_exog.iloc[:train_end]  
        test_exog = df_exog.iloc[train_end:]
```

```
mse = evaluate_sarima_exog(item, train, test, test_size, best_params, m, train_exog, test_exog)
mse_all_items[item] = mse

return mse_all_items
```

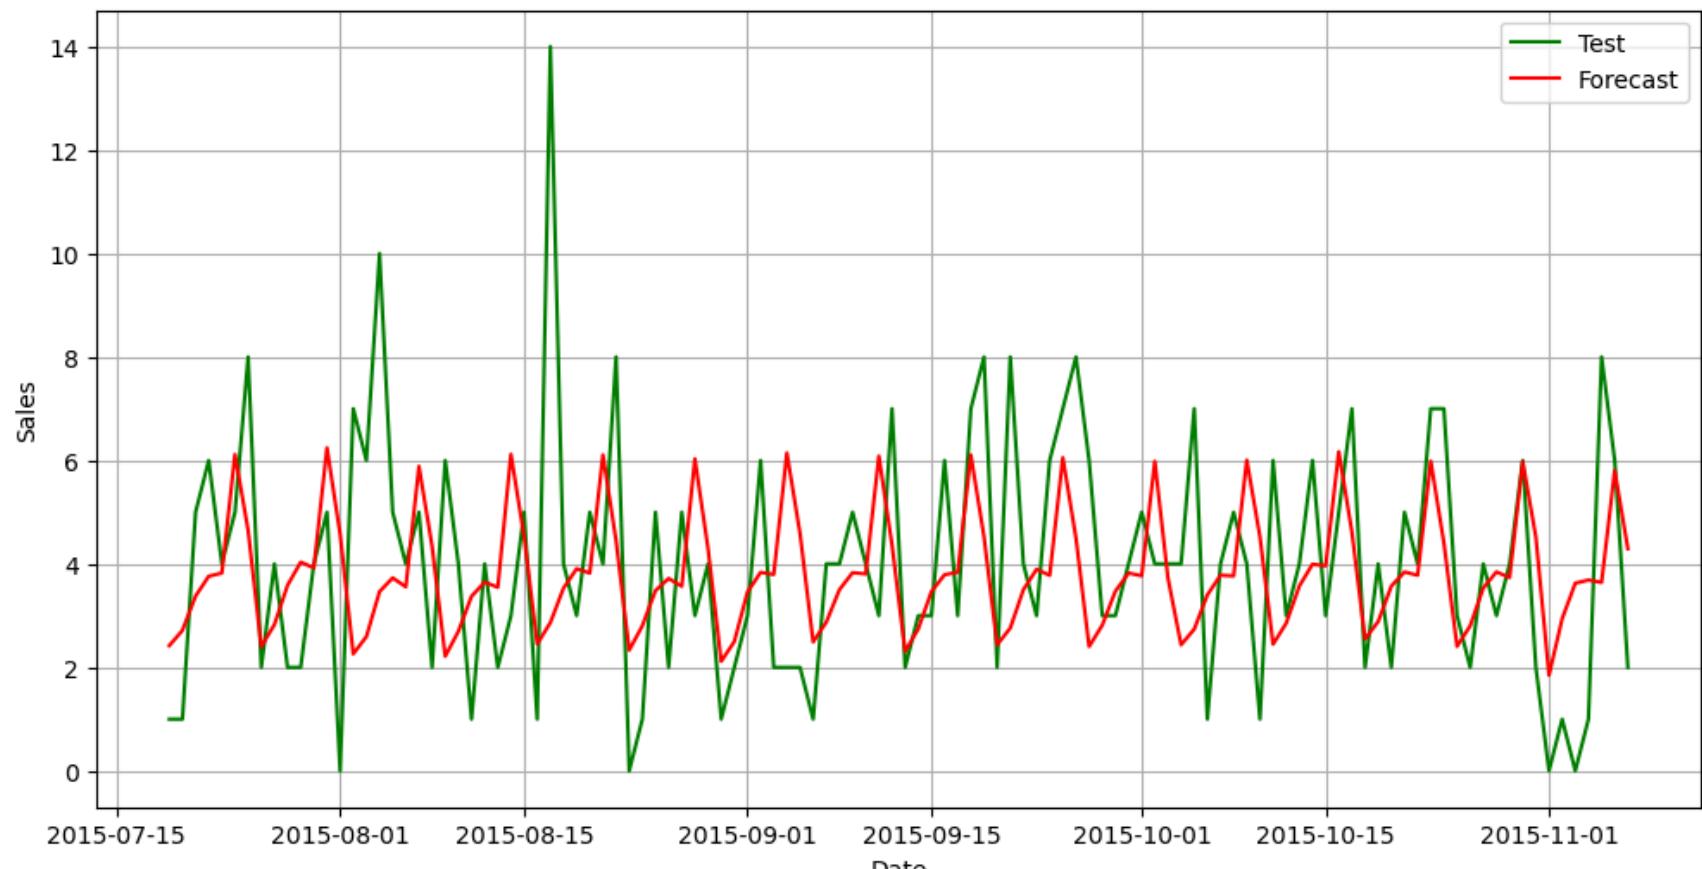
In [507...]

run_exog(7)



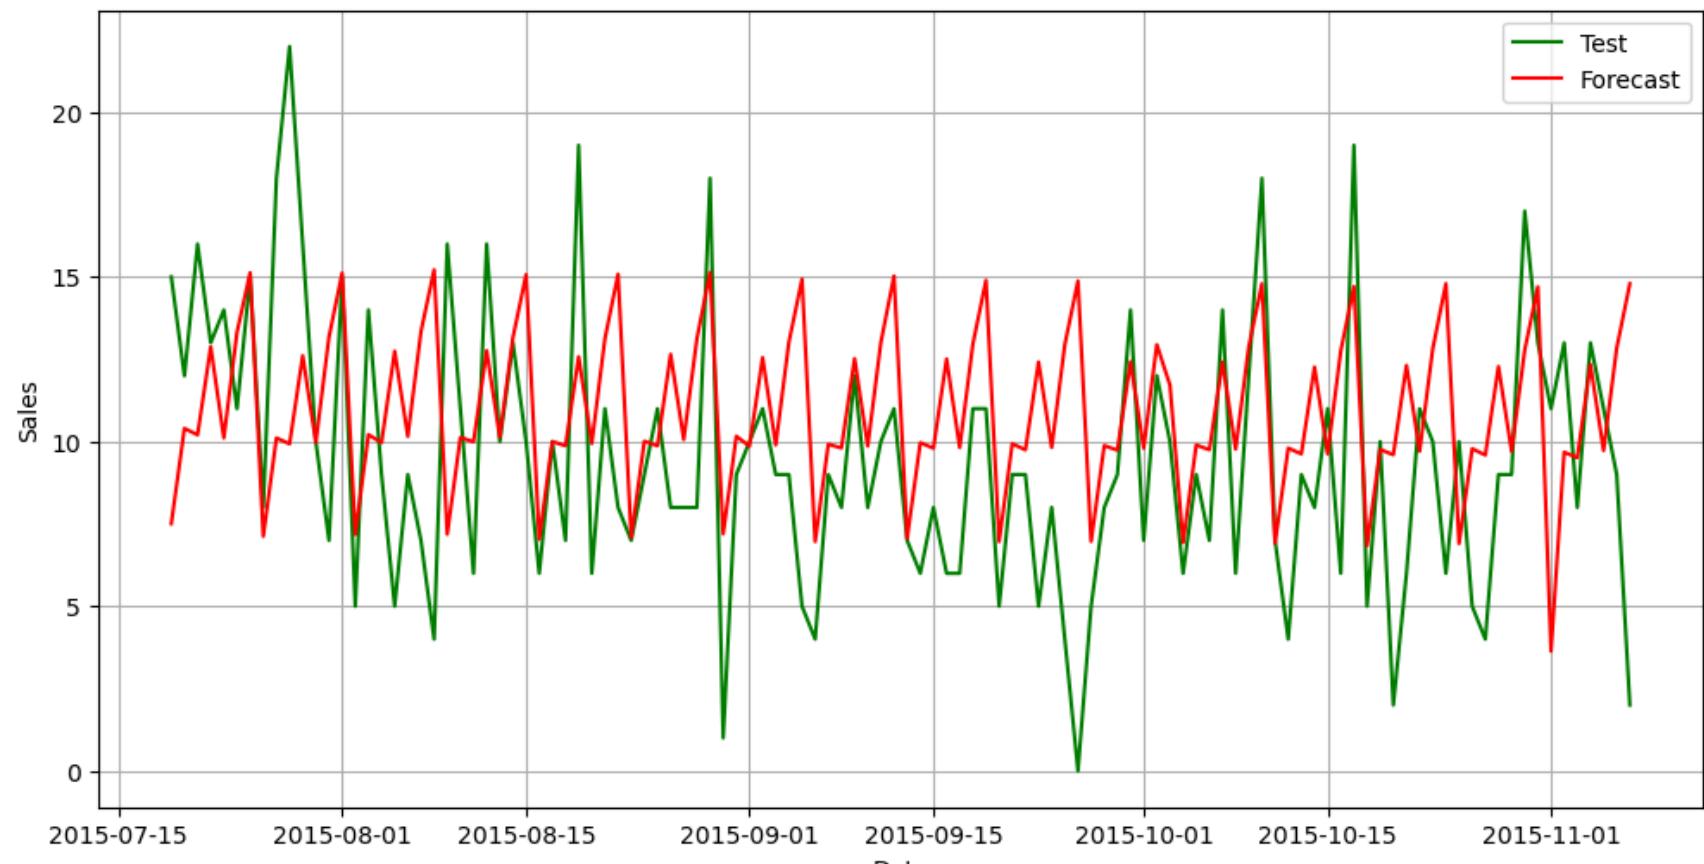
CALAMARI Mean Squared Error: 3.56

FISCH: SARIMA Forecast vs. Actual



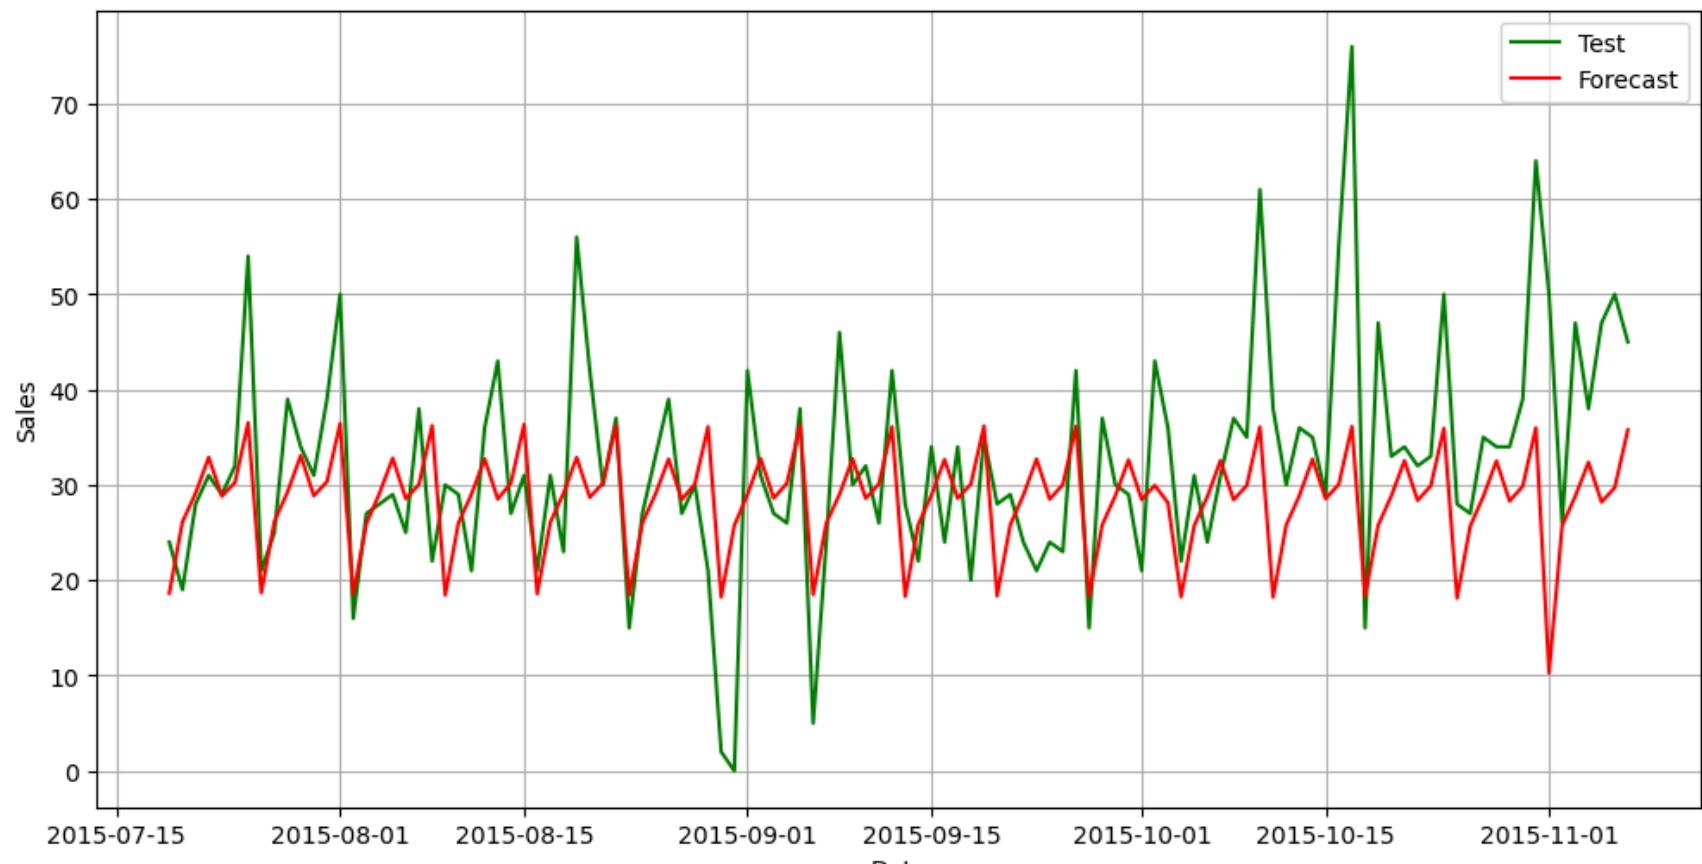
FISCH Mean Squared Error: 5.6

GARNELEN: SARIMA Forecast vs. Actual



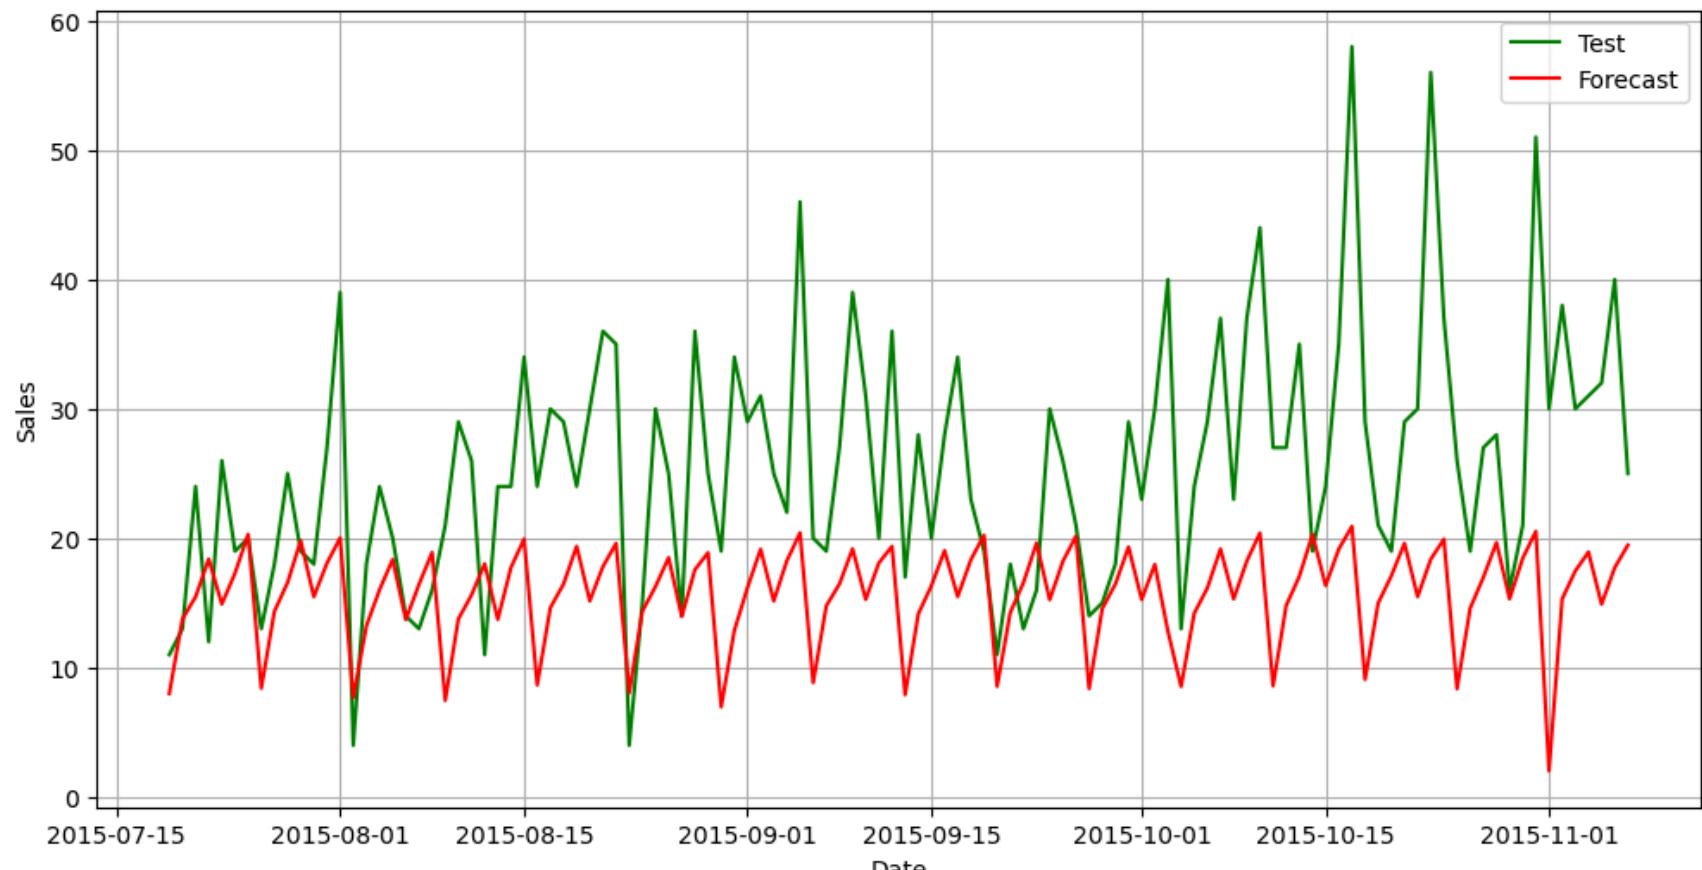
GARNELEN Mean Squared Error: 20.55

HAEHNCHEN: SARIMA Forecast vs. Actual



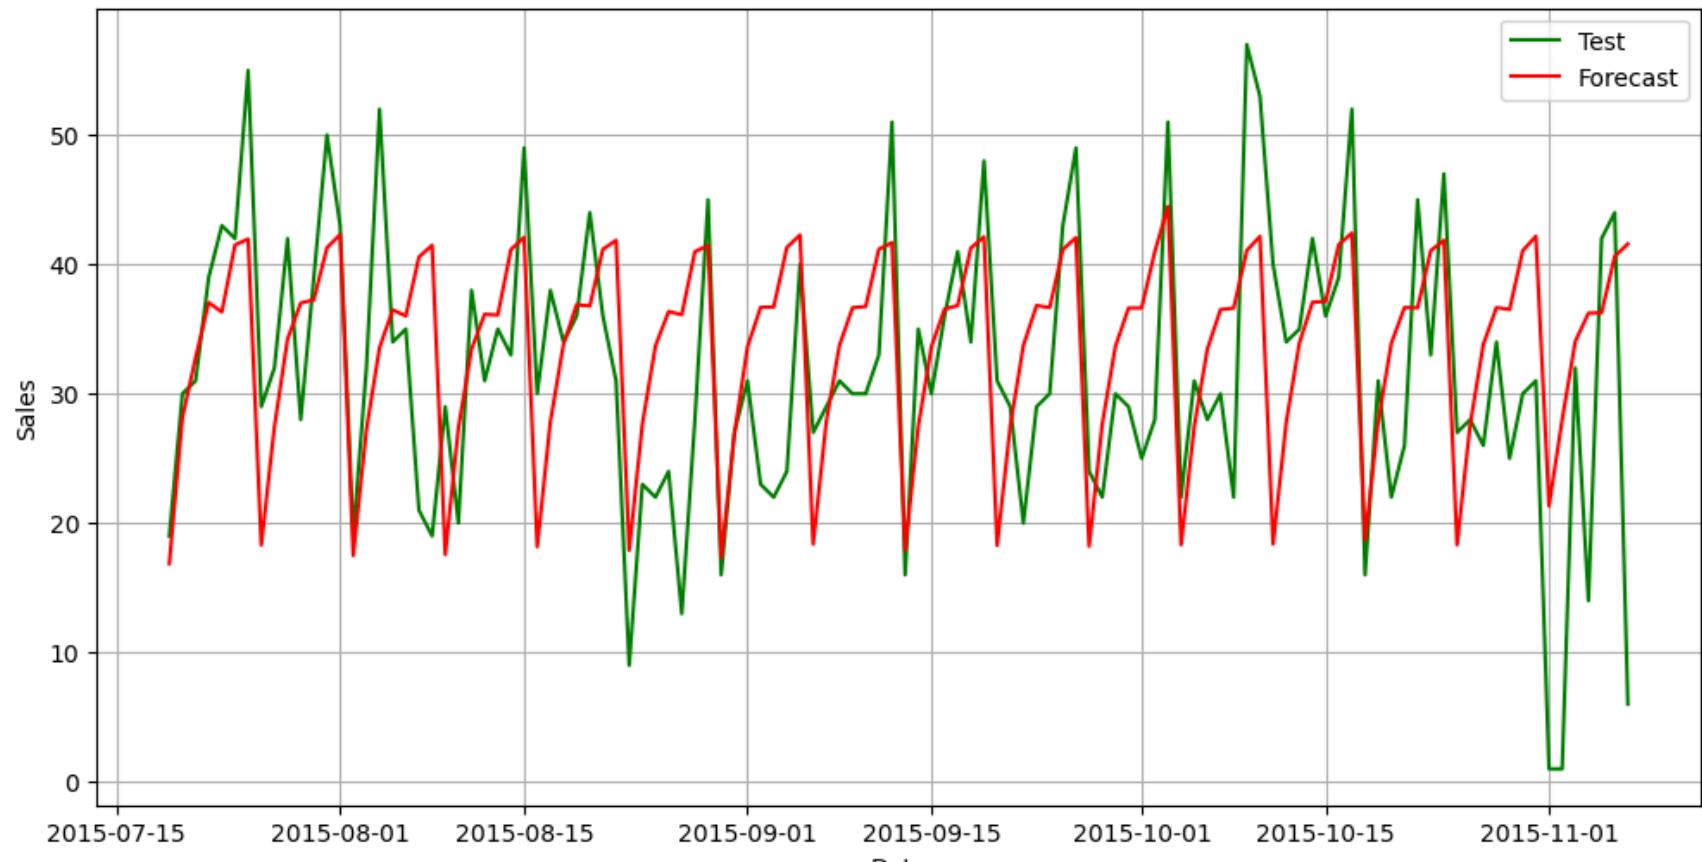
HAEHNCHEN Mean Squared Error: 119.78

KOEFT: SARIMA Forecast vs. Actual



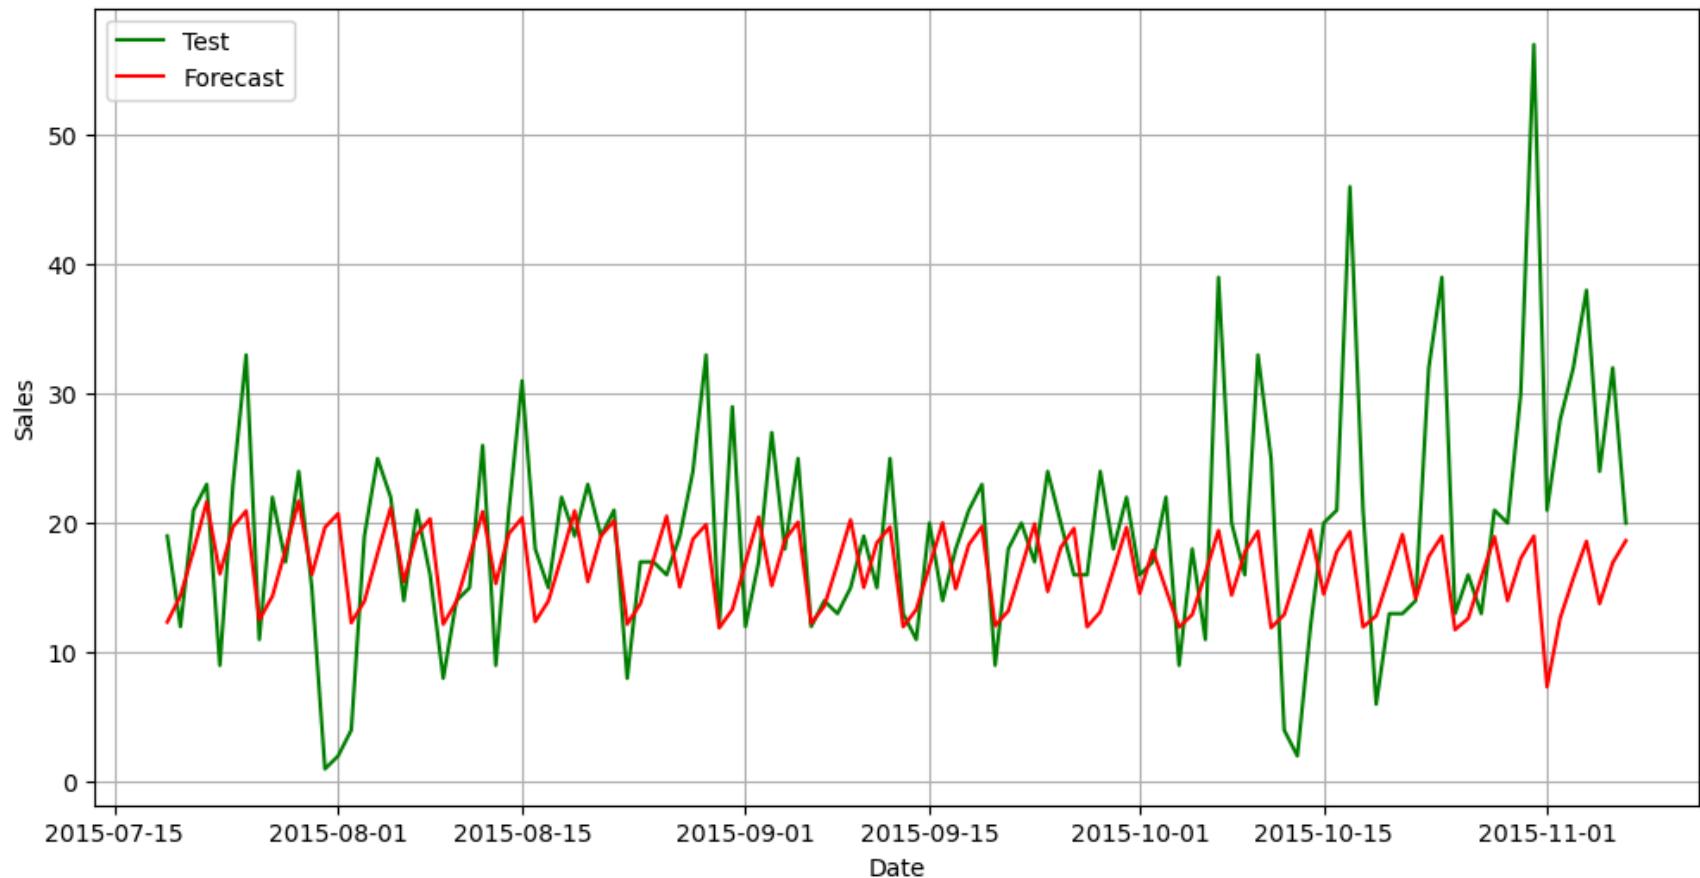
KOEFT Mean Squared Error: 171.68

LAMM: SARIMA Forecast vs. Actual



LAMM Mean Squared Error: 99.79

STEAK: SARIMA Forecast vs. Actual



STEAK Mean Squared Error: 76.28

Out[507]:

```
{'CALAMARI': 3.56,
 'FISCH': 5.6,
 'GARNELEN': 20.55,
 'HAEHNCHEN': 119.78,
 'KOEFTETE': 171.68,
 'LAMM': 99.79,
 'STEAK': 76.28}
```

In []: